

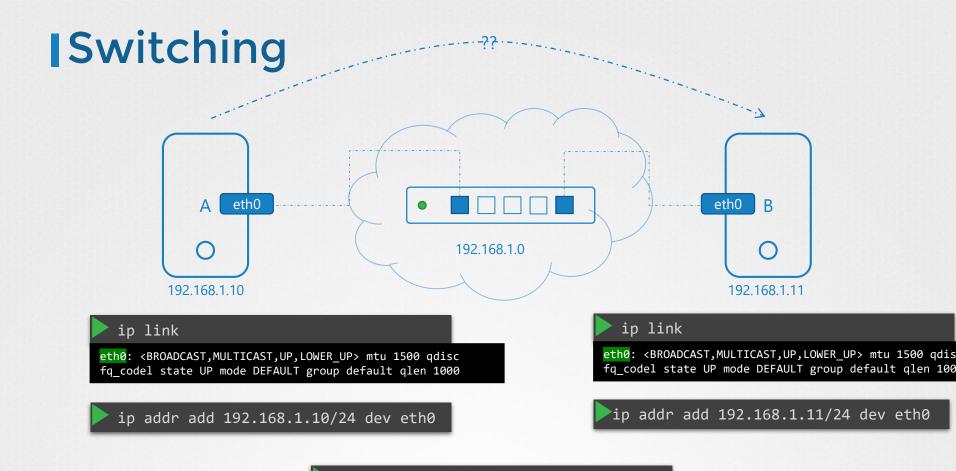


LINUX NETWORKING BASICS

| Networking Pre-Requisites

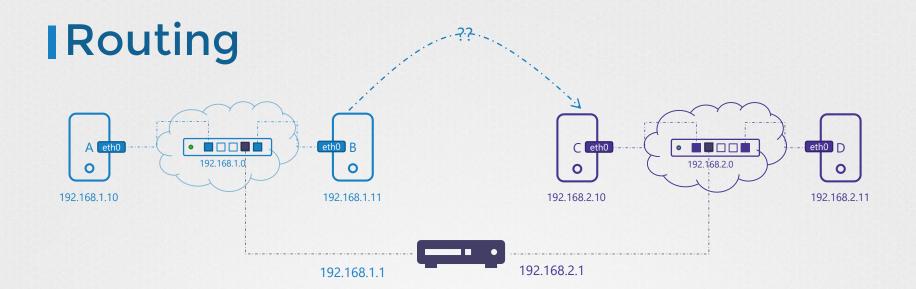
- Switching and Routing
 - Switching
 - Routing
 - Default Gateway
 - NAT
- Linux Interfaces for Virtual Networking
 - Bridge Network
 - VLAN
 - VXLAN
- IP Address Management & Name Resolution
 - DNS
 - IPAM
 - DHCP
- Firewalls
- Load-Balancers

- Tools:
 - Ping
 - NC NetCat
 - TCPDUMP
 - IPTABLES

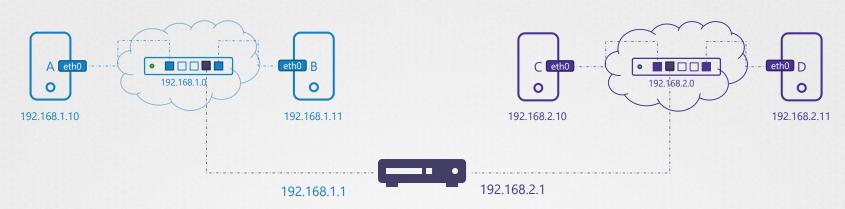


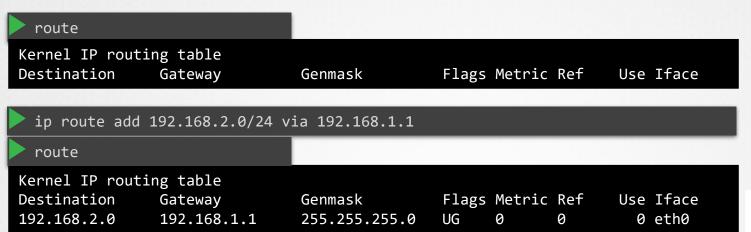
ping 192.168.1.11

Reply from 192.168.1.11: bytes=32 time=4ms TTL=117 Reply from 192.168.1.11: bytes=32 time=4ms TTL=117

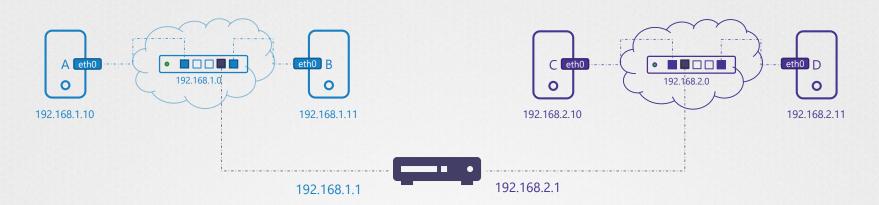


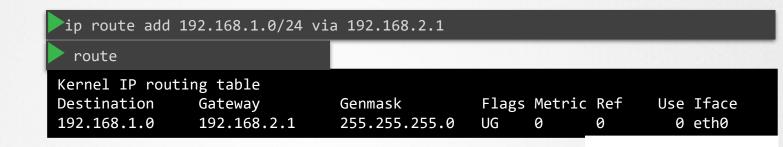
| Gateway

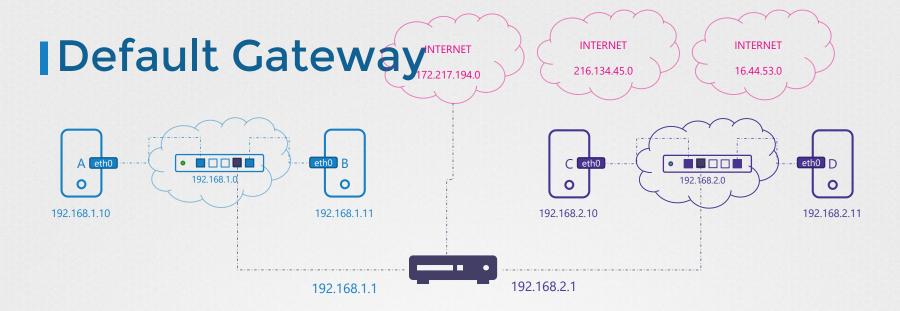




| Gateway





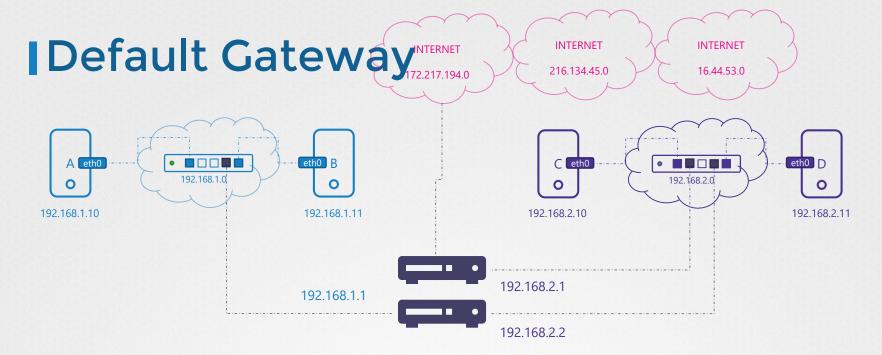


ip route add 192.168.1.0/24 via 192.168.2.1

ip route add default via 192.168.2.1

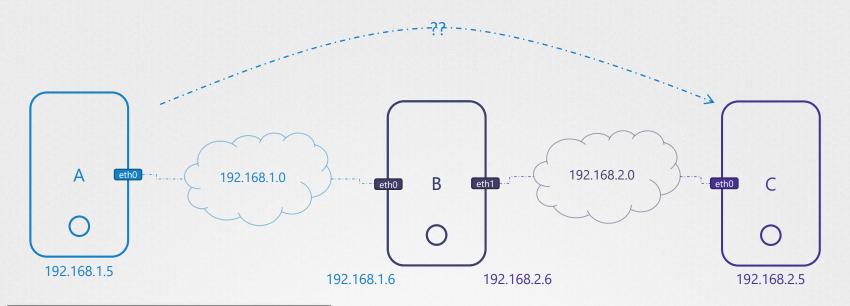
route

| Kernel IP routing table | | | | | | | |
|-------------------------|-------------|---------------|-------|--------|-----|-----|-------|
| Destination | Gateway | Genmask | Flags | Metric | Ref | Use | Iface |
| 192.168.1.0 | 192.168.2.1 | 255.255.255.0 | UG | 0 | 0 | 0 | eth0 |
| 0.0.0.0 | 192.168.2.1 | 255.255.255.0 | UG | 0 | 0 | 0 | eth0 |
| 192.168.2.0 | 0.0.0.0 | 255.255.255.0 | UG | 0 | 0 | 0 | eth0 |



ip route add 192.168.1.0/24 via 192.168.2.2

| П | Touce | | | | | | |
|---|-------------------------|-------------|---------------|-------|--------|-----|-----------|
| | Kernel IP routing table | | | | | | |
| | Destination | Gateway | Genmask | Flags | Metric | Ref | Use Iface |
| | default | 192.168.2.1 | 255.255.255.0 | UG | 0 | 0 | 0 eth0 |
| | 192.168.1.0 | 192.168.2.2 | 255.255.255.0 | UG | 0 | 0 | 0 eth0 |



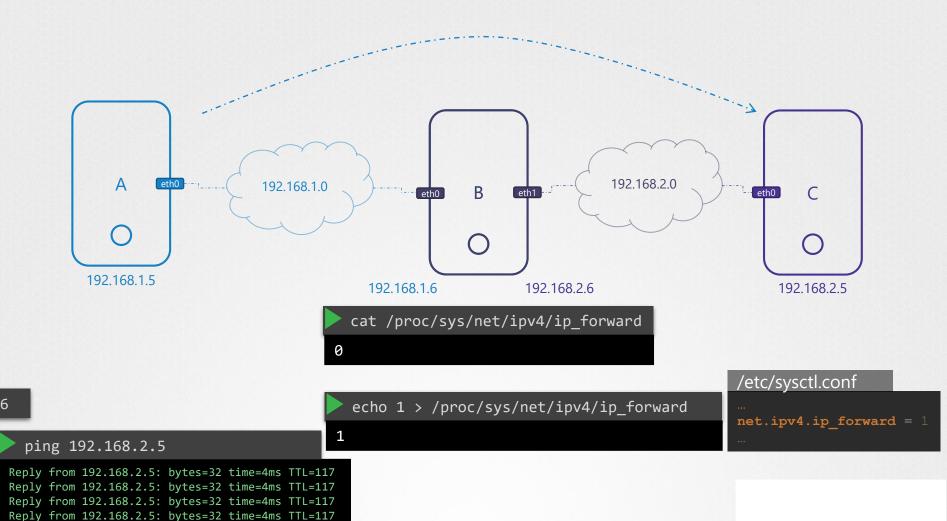
ping 192.168.2.5

Connect: Network is unreachable

ip route add 192.168.2.0/24 via <u>192.168.1.6</u>

ip route add 192.168.1.0/24 via 192.168.2.6

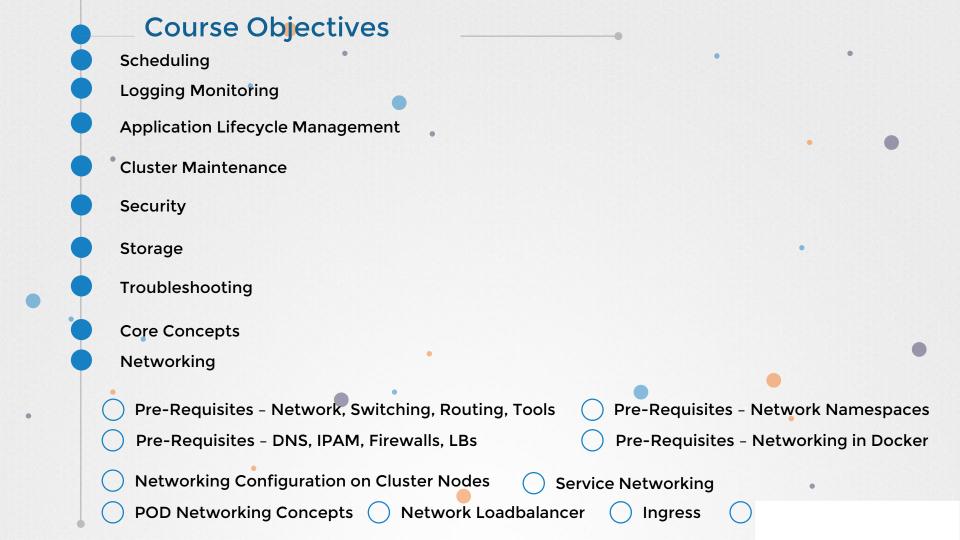
ping 192.168.2.5



Take Aways

```
ip link
 ip addr
 ip addr add 192.168.1.10/24 dev eth0
 ip route
ip route add 192.168.1.0/24 via 192.168.2.1
 cat /proc/sys/net/ipv4/ip_forward
```

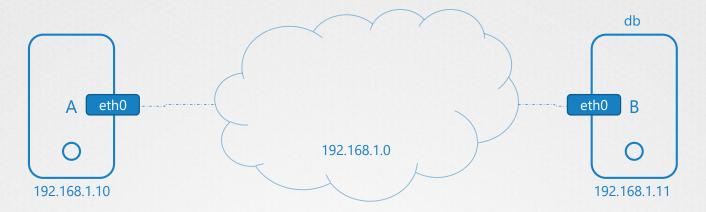
route





DNS

For the Absolute Beginners

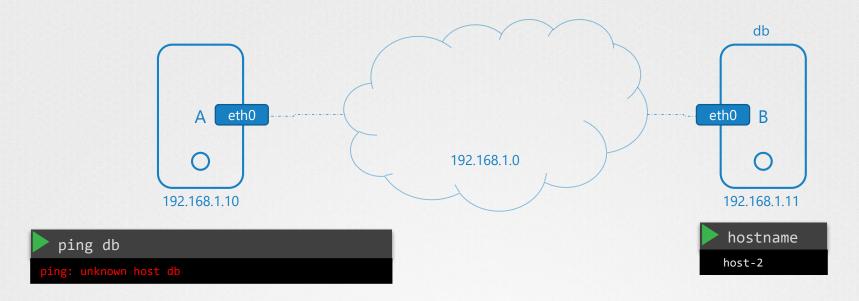


ping 192.168.1.11

Reply from 192.168.1.11: bytes=32 time=4ms TTL=117 Reply from 192.168.1.11: bytes=32 time=4ms TTL=117

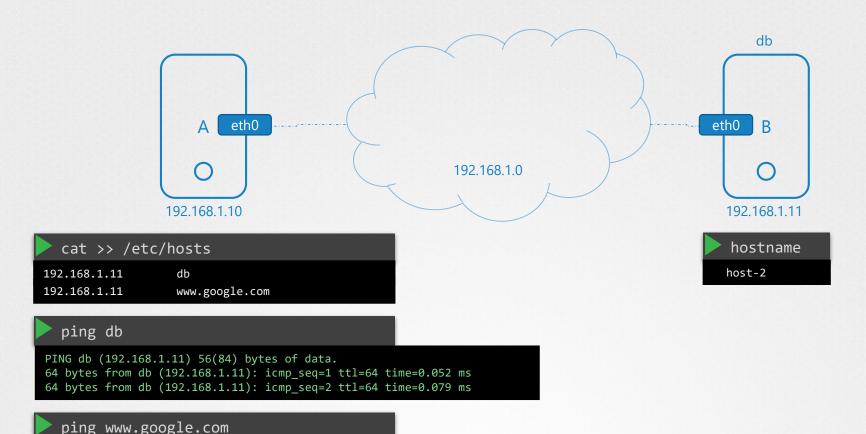
ping db

ping: unknown host db



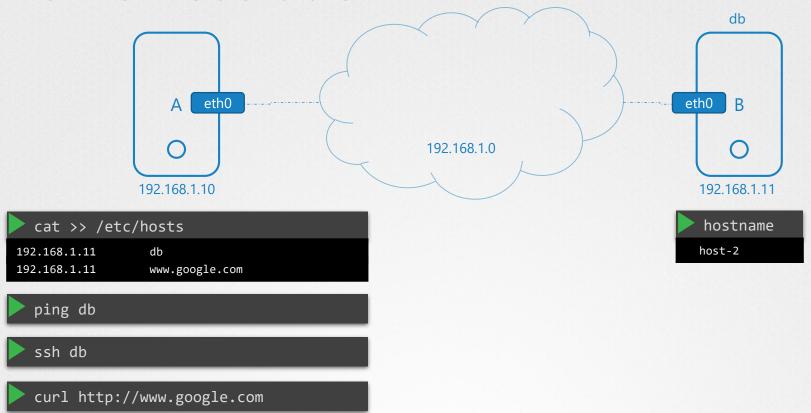


```
PING db (192.168.1.11) 56(84) bytes of data.
64 bytes from db (192.168.1.11): icmp_seq=1 ttl=64 time=0.052 ms
64 bytes from db (192.168.1.11): icmp_seq=2 ttl=64 time=0.079 ms
```

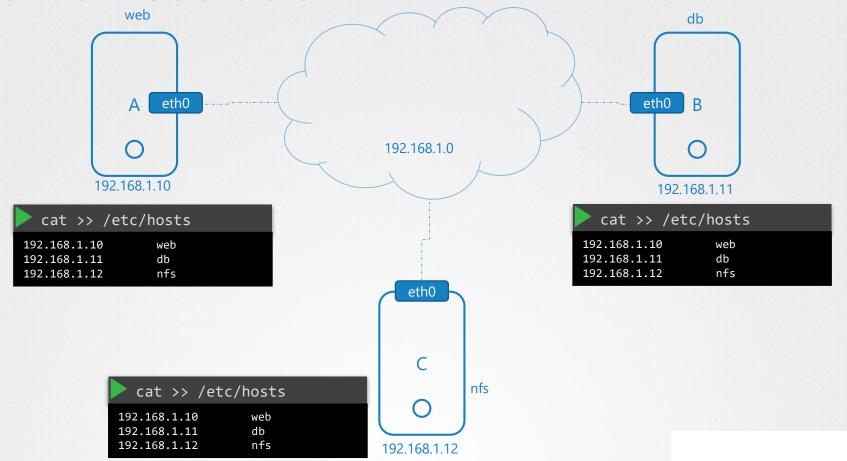


PING www.google.com (192.168.1.11) 56(84) bytes of data.
64 bytes from www.google.com (192.168.1.11): icmp_seq=1 ttl=64 time=0.052 ms
64 bytes from www.google.com (192.168.1.11): icmp_seq=2 ttl=64 time=0.079 ms

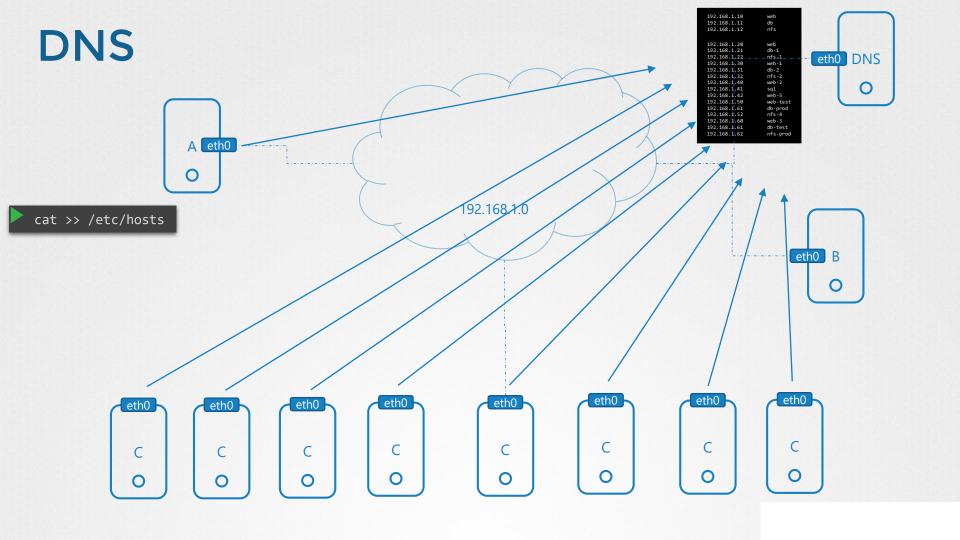
Name Resolution



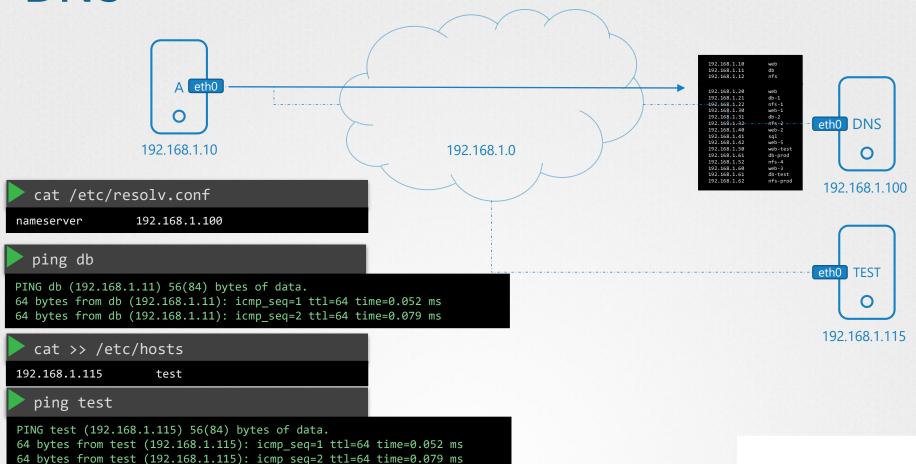
Name Resolution

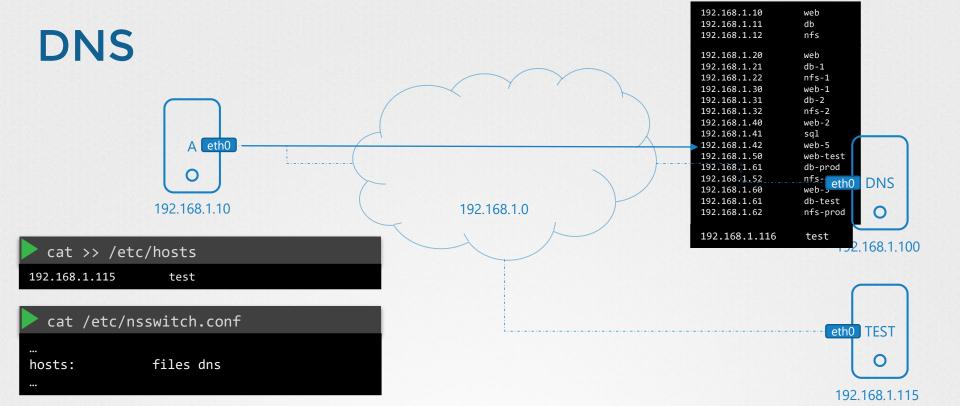


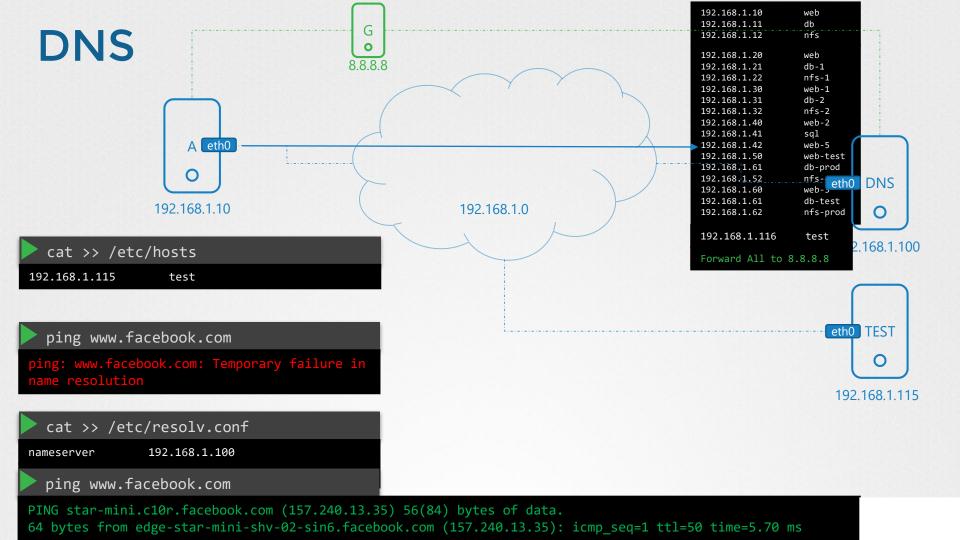
Name Resolution A eth0 eth0 B 0 0 192.168.1.0 cat >> /etc/hosts 192.168.1.10 web 192.168.1.11 db 192.168.1.12 nfs 192.168.1.20 web 192.168.1.21 db-1 192.168.1.22 nfs-1 192.168.1.30 web-1 192.168.1.31 db-2 192.168.1.32 nfs-2 192.168.1.40 web-2 192.168.1.41 sql 192.168.1.42 web-5 192.168.1.50 web-test 192.168.1.61 db-prod 192.168.1.52 nfs-4 eth0 eth0 eth0 eth0 eth0 192.168.1.60 web-3 192.168.1.61 db-test 192.168.1.62 nfs-prod W 0



DNS







I Domain Names

wwwkubernetes.io

www.codepen.io

www.facebook.com

www.un.org

www.mit.edu

www.google.com

www.behance.net

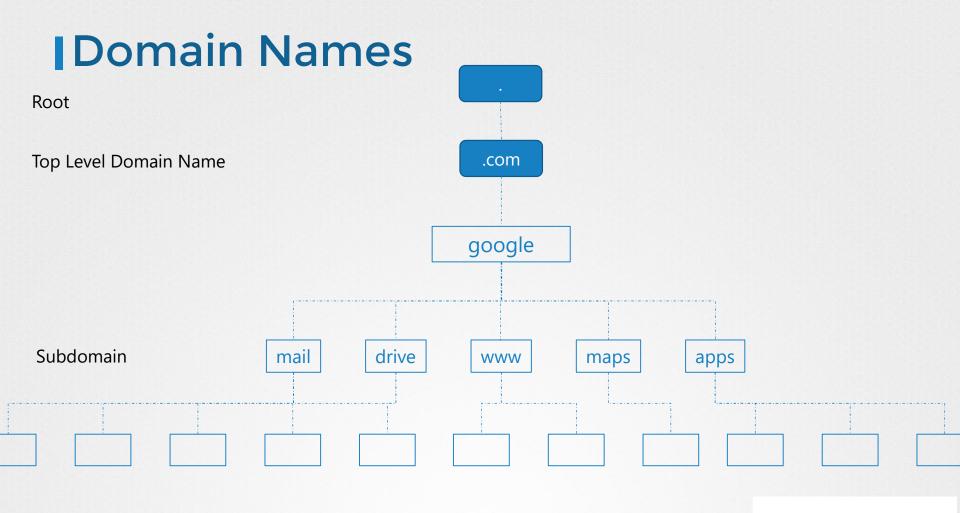
www.speedtest.net

www.stanford.edu

www.care.org

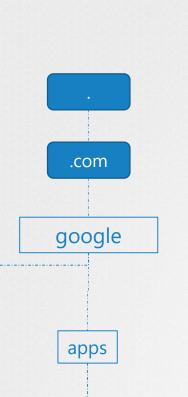
I Domain Names

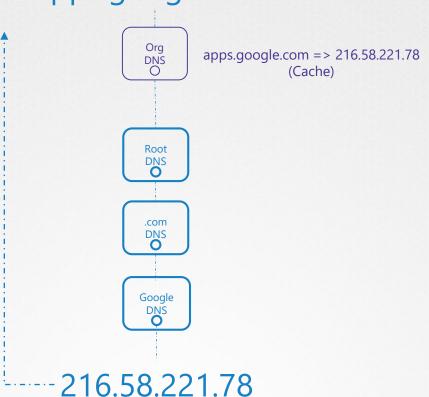




Domain Names

apps.google.com





| Domain Names



ISearch Domain



192.168.1.10 web.mycompany.com 192.168.1.11 db.mycompany.com 192,168,1,12 nfs.mycompany.com 192.168.1.13 web-1.mycompany.com 192.168.1.14 sql.mycompany.com

mycompany.com

nfs

web

mail

drive

WWW

pay

hr

sql

cat >> /etc/resolv.conf

nameserver

192.168.1.100

mycompany.com prod.mycompany.com search

ping web

PING web (192.168.1.10) 56(84) bytes of data.

64 bytes from web (192.168.1.10): icmp seq=1 ttl=64 time=0.052 ms 64 bytes from web (192.168.1.10): icmp seq=2 ttl=64 time=0.079 ms

ping web

PING web.mycompany.com (192.168.1.10) 56(84) bytes of data.

64 bytes from web.mycompany.com (192.168.1.10): ... time=0.052 ms

64 bytes from web.mycompany.com (192.168.1.10): ... time=0.079 ms

PING web.mycompany.com (192.168.1.10) 56(84) bytes of data.

ning web

ping web.mycompany.com

ping web.mycompany.com

PTNG web.mycompany.com (192.168.1.10) 56(84) bytes of data.

rom web.mycompany.com (192.168.1.10): ttl=64 time=0.052 ms

64 bytes from web.mycompany.com (192.168.1.10): ttl=64 time=0.052 ms

IRecord Types

| Α | web-server | 192.168.1.1 |
|-------|-----------------|-----------------------------------------|
| AAAA | web-server | 2001:0db8:85a3:0000:0000:8a2e:0370:7334 |
| CNAME | food.web-server | eat.web-server, hungry.web-server |

Inslookup

```
nslookup www.google.com

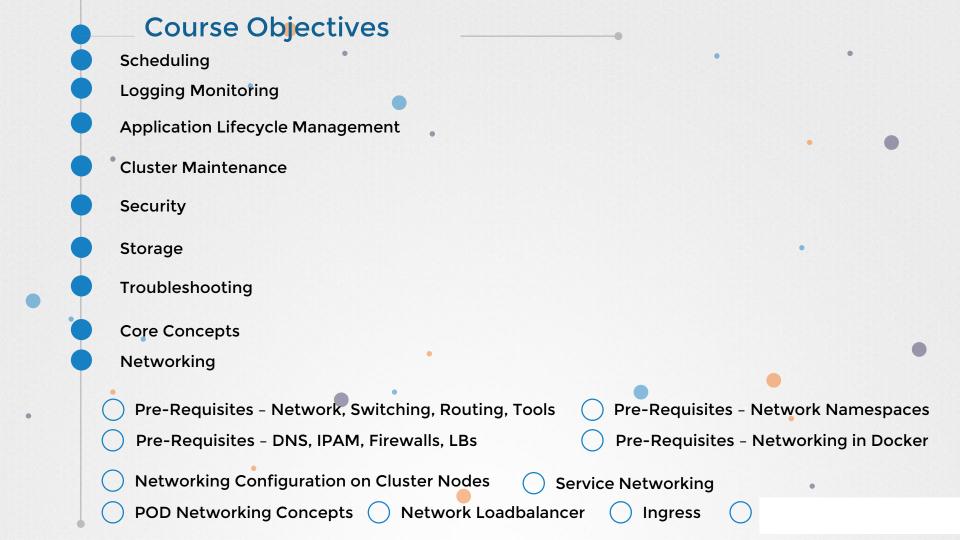
Server: 8.8.8.8
Address: 8.8.8.8#53

Non-authoritative answer:
Name: www.google.com
Address: 172.217.0.132
```

Idig

dig www.google.com

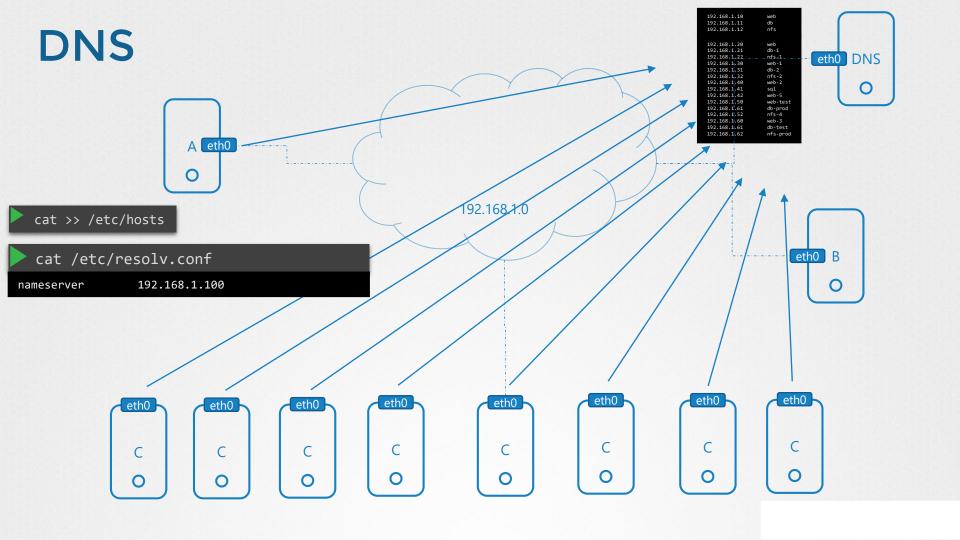
```
; <<>> DiG 9.10.3-P4-Ubuntu <<>> www.google.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 28065
;; flags: qr rd ra; QUERY: 1, ANSWER: 6, AUTHORITY: 0, ADDITIONAL: 1
;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 512
;; QUESTION SECTION:
;www.google.com.
                                        ΙN
;; ANSWER SECTION:
www.google.com.
                        245
                                IN
                                                64.233.177.103
www.google.com.
                        245
                                IN
                                                64.233.177.105
www.google.com.
                        245
                                IN
                                                64.233.177.147
www.google.com.
                        245
                                IN
                                                64.233.177.106
www.google.com.
                        245
                                IN
                                                64.233.177.104
www.google.com.
                                                64.233.177.99
                        245
                                IN
;; Query time: 5 msec
;; SERVER: 8.8.8.8#53(8.8.8.8)
;; WHEN: Sun Mar 24 04:34:33 UTC 2019
;; MSG SIZE rcvd: 139
```



CoreDNS



Name Resolution A eth0 eth0 B 0 0 192.168.1.0 cat >> /etc/hosts 192.168.1.10 web 192.168.1.11 db 192.168.1.12 nfs 192.168.1.20 web 192.168.1.21 db-1 192.168.1.22 nfs-1 192.168.1.30 web-1 192.168.1.31 db-2 192.168.1.32 nfs-2 192.168.1.40 web-2 192.168.1.41 sql 192.168.1.42 web-5 192.168.1.50 web-test 192.168.1.61 db-prod 192.168.1.52 nfs-4 eth0 eth0 eth0 eth0 eth0 192.168.1.60 web-3 192.168.1.61 db-test 192.168.1.62 nfs-prod W 0



CoreDNS



| web db |
|------------------------------------------------------------------|
| web db-1 nfs-1 web-1 db-2 nfs-2 web-2 sq1 web-5 web-test db-prod |
| nfs-4 web-3 db-test nfs-prod |
| |





wget https://github.com/coredns/coredns/releases/download/v1.4.0/coredns_1.4.0_linux_amd64.tgz

coredns_1.4.0_linux_amd64.tgz

tar -xzvf coredns_1.4.0_linux_amd64.tgz

./coredns .:53 2019-03-04T10:46:13.756Z [INFO] CoreDNS-1.4.0 2019-03-04T10:46:13.756Z [INFO] linux/amd64, go1.12, 8dcc7fc CoreDNS-1.4.0 linux/amd64, go1.12, 8dcc7fc

| 192.168.1.10 192.168.1.11 | web db |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------|
| 192.168.1.20 192.168.1.21 192.168.1.22 192.168.1.30 192.168.1.31 192.168.1.32 192.168.1.44 192.168.1.44 192.168.1.42 192.168.1.50 192.168.1.50 192.168.1.50 192.168.1.50 | web db-1 nfs-1 web-1 db-2 nfs-2 web-2 sql web-5 web-test db-prod nfs-4 web-3 db-test |
| 192.168.1.62 | nfs-prod |





cat /etc/hosts

| 192.168.1.10 192.168.1.11 | web db |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------|
| 192.168.1.20 192.168.1.21 192.168.1.22 192.168.1.30 192.168.1.31 192.168.1.32 192.168.1.40 192.168.1.41 192.168.1.42 192.168.1.50 192.168.1.61 | web db-1 nfs-1 web-1 db-2 nfs-2 web-2 sql web-5 web-test db-prod |
| 192.168.1.52 192.168.1.60 192.168.1.61 192.168.1.62 | nfs-4 web-3 db-test nfs-prod |
| | |



cat > /etc/hosts 192.168.1.10 web 192.168.1.11 db 192.168.1.20 web 192.168.1.21 db-1 192.168.1.22 nfs-1 192.168.1.30 web-1 db-2 192.168.1.31 192.168.1.32 nfs-2 192.168.1.40 web-2 192.168.1.41 sql 192.168.1.42 web-5 192.168.1.50 web-test 192.168.1.61 db-prod 192.168.1.52 nfs-4 192.168.1.60 web-3 192.168.1.61 db-test nfs-prod 192.168.1.62

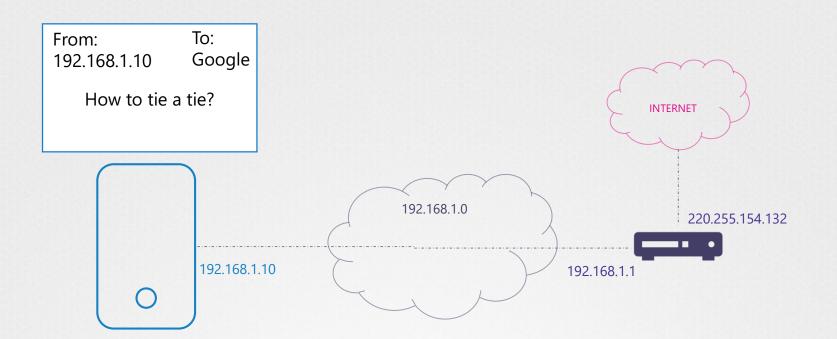
```
cat > Corefile
. {
    hosts /etc/hosts
}
```

```
./coredns
.:53
2019-03-04T10:46:13.756Z [INFO] CoreDNS-1.4.0
2019-03-04T10:46:13.756Z [INFO] linux/amd64, go1.12, 8dcc7fc
CoreDNS-1.4.0
linux/amd64, go1.12, 8dcc7fc
```





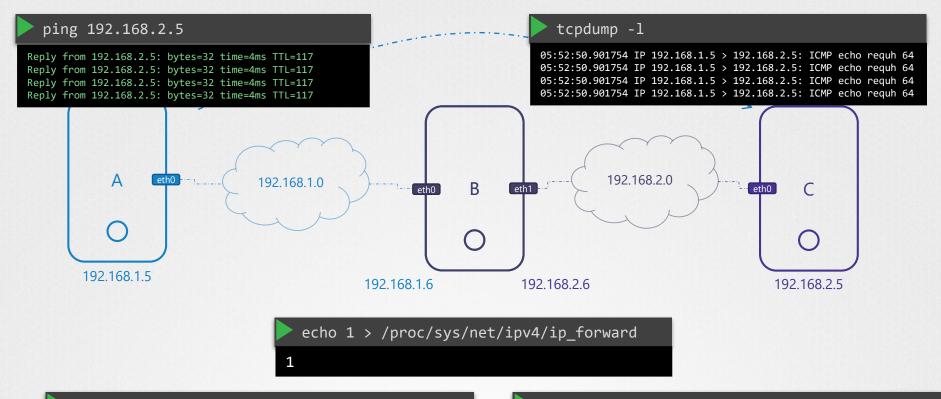
PRE-REQUISITE NETWORK ADDRESS TRANSLATION (NAT)



Network Address Translation (NAT) To: From: 220.255.154.132 Google Here's how.... 192.168.1.0 .154.132 From: To: From: From: 192.168.1.10 220.255.154.132 192.168.1.10 Google How to tie a tie? 92.168.1.10 192.168.1.1

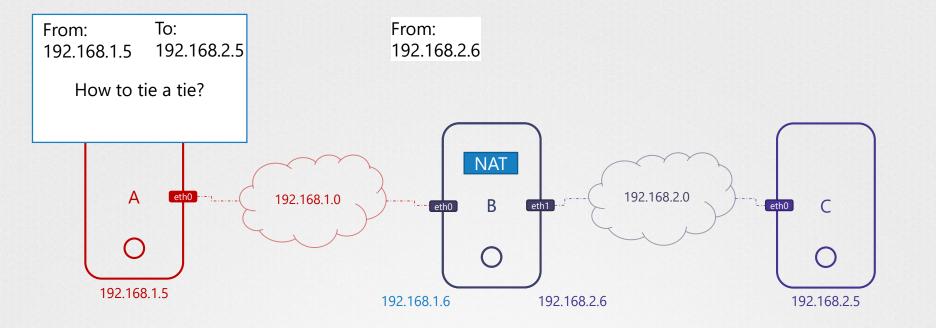
NAT TABLE

| Private IP | Public IP |
|--------------|-----------------|
| 192.168.1.10 | 220.255.154.132 |
| | |



ip route add 192.168.2.0/24 via 192.168.1.6

ip route add 192.168.1.0/24 via 192.168.2.6

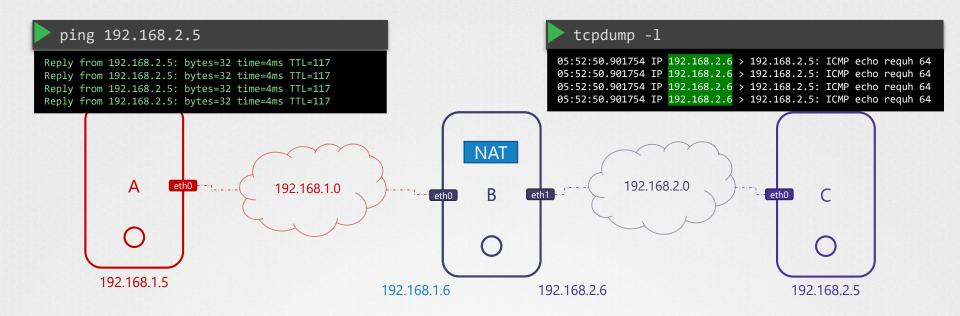






ip route add 192.168.2.0/24 via 192.168.1.6

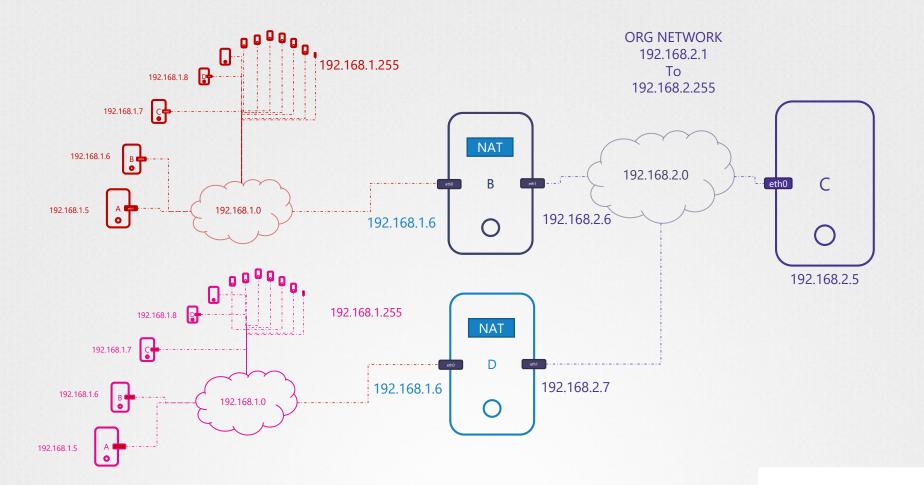
ip route add 192.168.1.0/24 via 192.168.2.6



iptables -t nat -A POSTROUTING -s 192.168.5.0/24 -j MASQUERADE

echo 1 > /proc/sys/net/ipv4/ip_forward
1

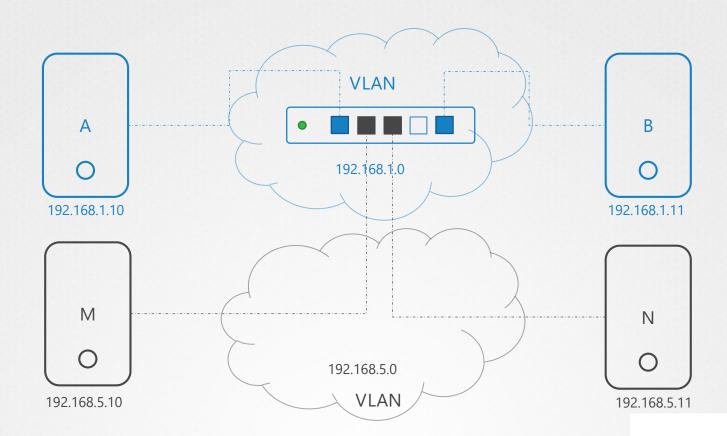
ip route add 192.168.2.0/24 via 192.168.1.6





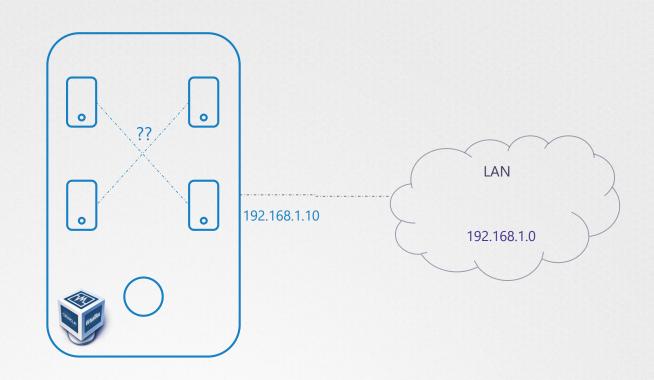
NETWORKING VLAN & VXLAN

IVLAN

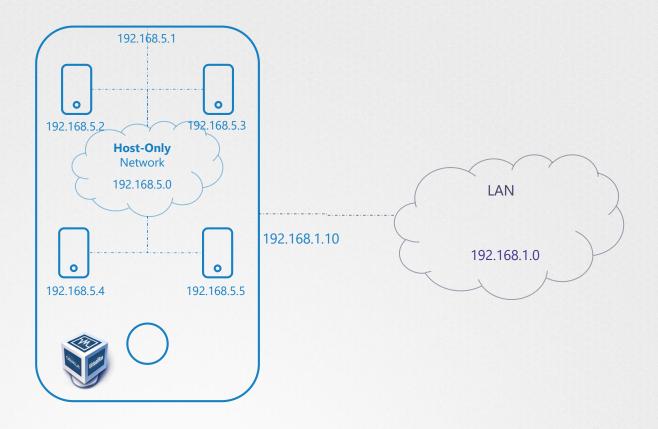




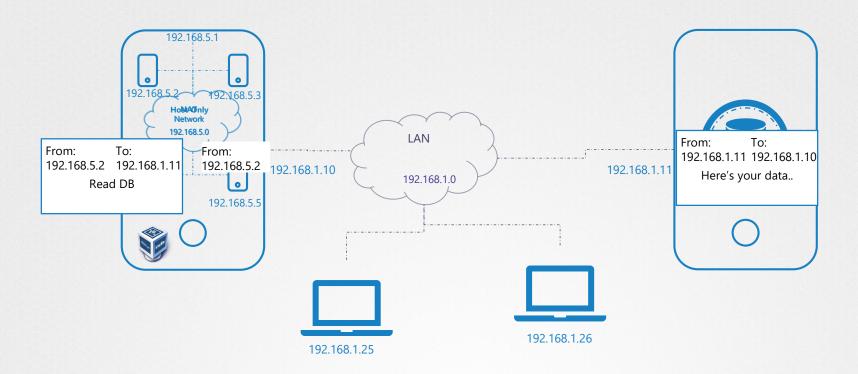
NETWORKING VIRTUAL MACHINES



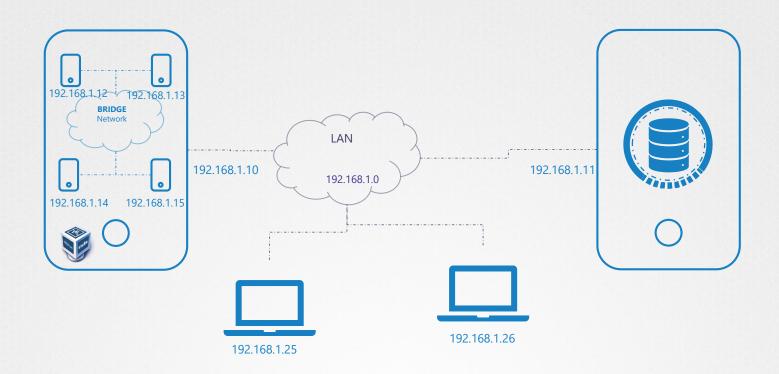
Host Only



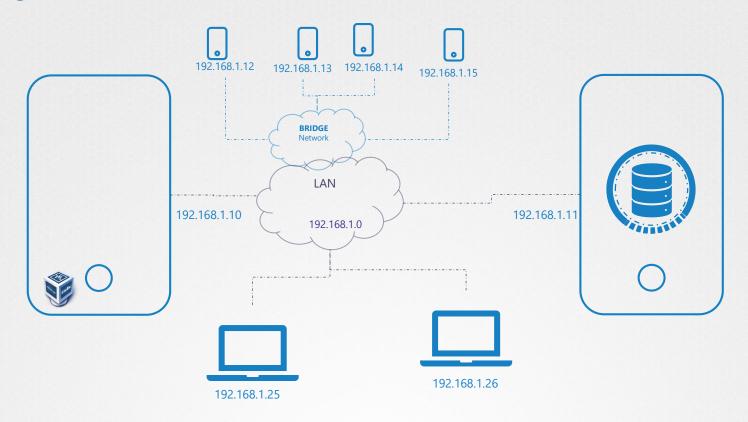
Network Address Translation (NAT)



Bridge Network



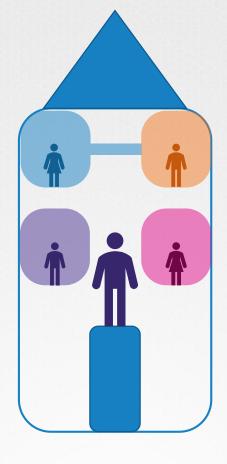
Bridge Network



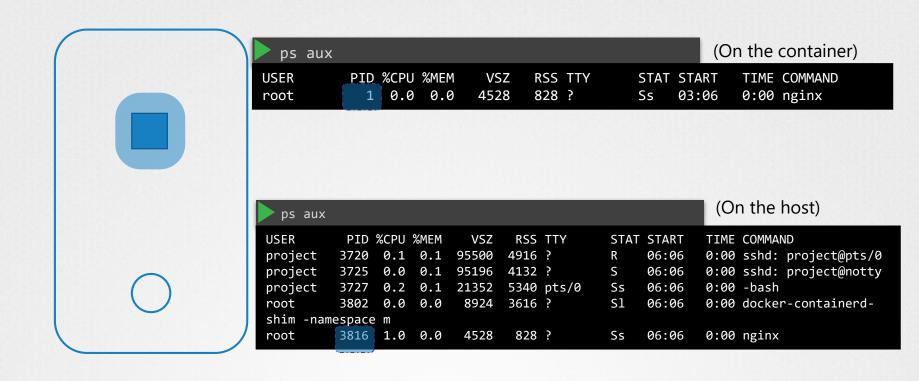


NETWORK
NAMESPACES

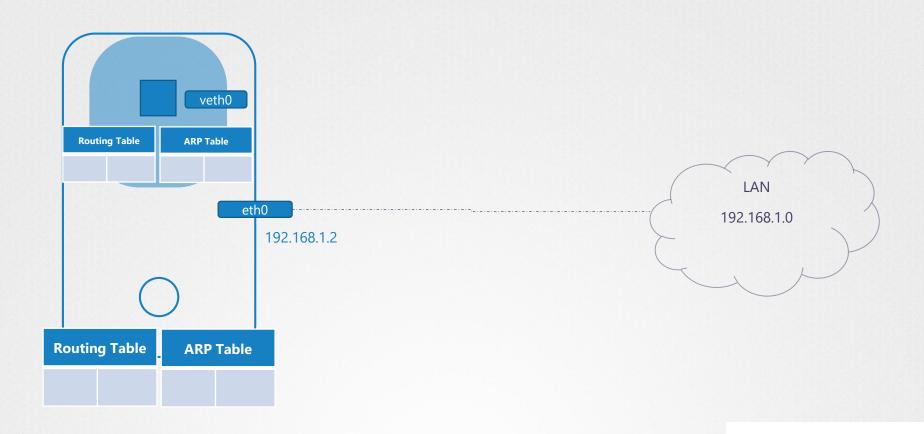
NAMESPACE



PROCESS NAMESPACE

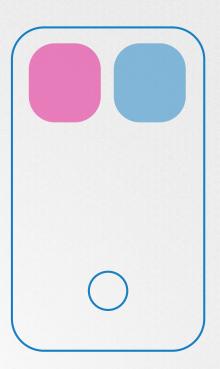


NETWORK NAMESPACE



CREATE NETWORK NS





EXEC IN NETWORK NS

ip link

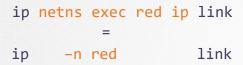
- 1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc state UNKNOWN mode DEFAULT group default qlen 1000 link/loopback 00:00:00:00:00:00:00:00:00:00:00:00
- 2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc state UP mode DEFAULT qlen 1000 link/ether 02:42:ac:11:00:08 brd ff:ff:ff:ff:ff

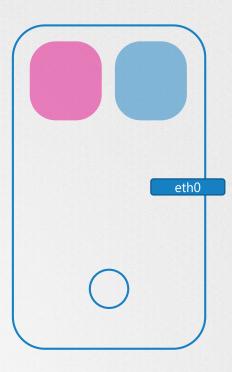
ip hėnks exec red

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc state UNKNOWN mode DEFAULT group default qlen 1000 link/loopback 00:00:00:00:00:00:00:00:00:00:00:00

ip linked

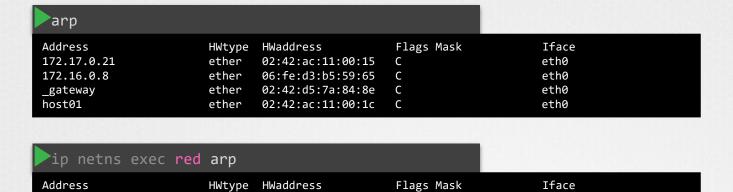
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc state UNKNOWN mode DEFAULT group default qlen 1000 link/loopback 00:00:00:00:00:00:00:00:00:00:00:00:00

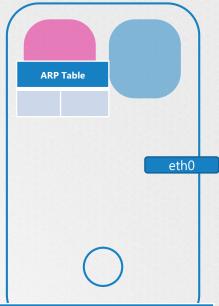




EXEC IN NETWORK NS

HWtype HWaddress





| ARP Table | | |
|-------------|-------------------|--|
| 172.17.0.21 | 02:42:ac:11:00:15 | |
| 172.16.0.8 | 06:fe:d3:b5:59:65 | |

EXEC IN NETWORK NS

route

| Kernel IP rout | ing table | | | | | | |
|----------------|-----------|---------------|-------|--------|-----|-----|---------|
| Destination | Gateway | Genmask | Flags | Metric | Ref | Use | Iface |
| default | _gateway | 0.0.0.0 | UG | 202 | 0 | 0 | eth0 |
| 172.17.0.0 | 0.0.0.0 | 255.255.0.0 | U | 202 | 0 | 0 | eth0 |
| 172.17.0.0 | 0.0.0.0 | 255.255.255.0 | U | 0 | 0 | 0 | docker@ |

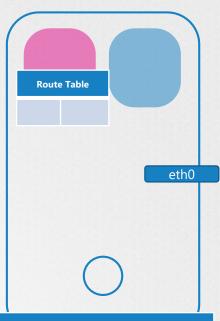
ip netns exec red route

Kernel IP routing table
Destination Gateway

teway Genmask

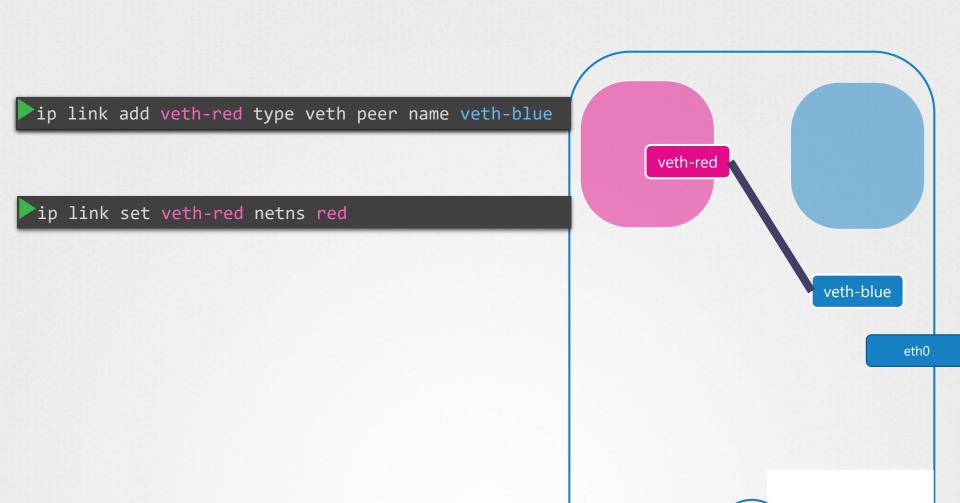
Flags Metric Ref

Use Iface



| Route Table | | |
|-------------|---------|--|
| 172.17.0.0 | 0.0.0.0 | |
| 17.18.0.0 | 0.0.0.0 | |

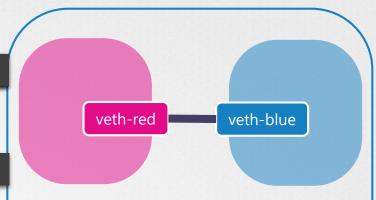
ip link add veth-red type veth peer name veth-blue veth-red veth-blue eth0





ip link set veth-red netns red

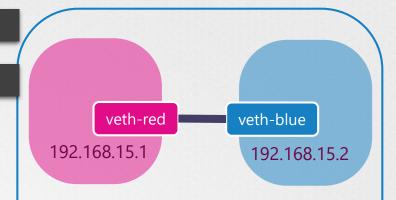
ip link set veth-blue netns blue



eth0

- ▶ip link add veth-red type veth peer name veth-blue
- ip link set veth-red netns red
- ip link set veth-blue netns blue

- ip -n red addr add 192.168.15.1 dev veth-red
- ip -n blue addr add 192.168.15.2 dev veth-blue

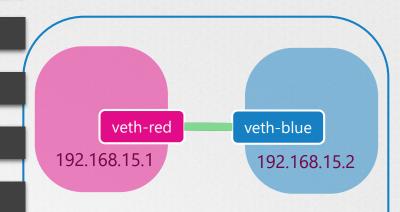


eth0

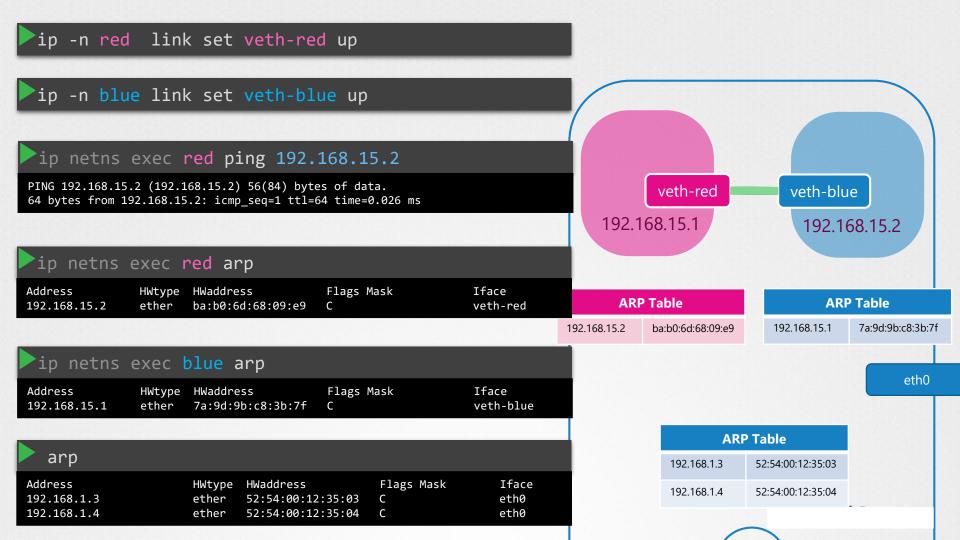
- ip link add veth-red type veth peer name veth-blue
- ip link set veth-red netns red
- ip link set veth-blue netns blue
- ip -n red addr add 192.168.15.1 dev veth-red
- ip -n blue addr add 192.168.15.2 dev veth-blue

- ▶ip -n red link set veth-red up
- ip -n blue link set veth-blue up
- ip netns exec red ping 192.168.15.2

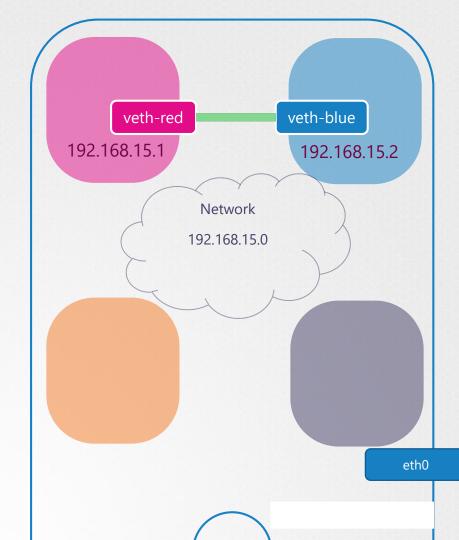
PING 192.168.15.2 (192.168.15.2) 56(84) bytes of data. 64 bytes from 192.168.15.2: icmp_seq=1 ttl=64 time=0.026 ms



eth0





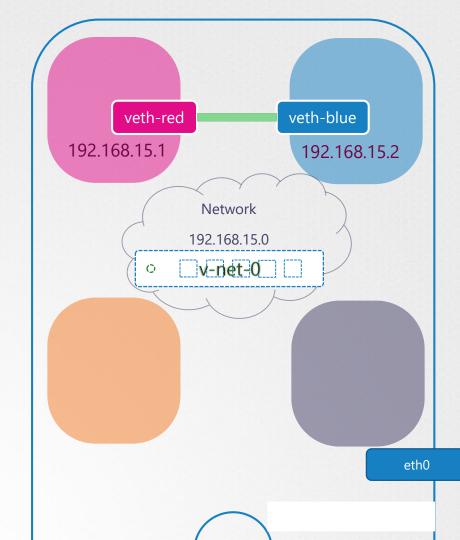


ip link add v-net-0 type bridge

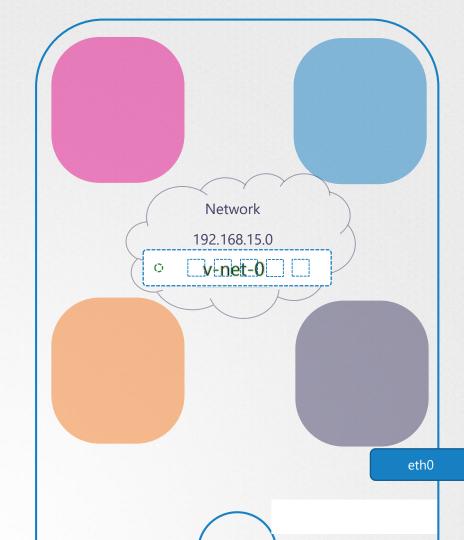
```
ip link
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state
UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc
fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 02:0d:31:14:c7:a7 brd ff:ff:ff:ff:ff
6: v-net-0: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN
mode DEFAULT group default qlen 1000
    link/ether 06:9d:69:52:6f:61 brd ff:ff:ff:ff:ff
```

ip link set dev v-net-0 up



veth-red veth-blue

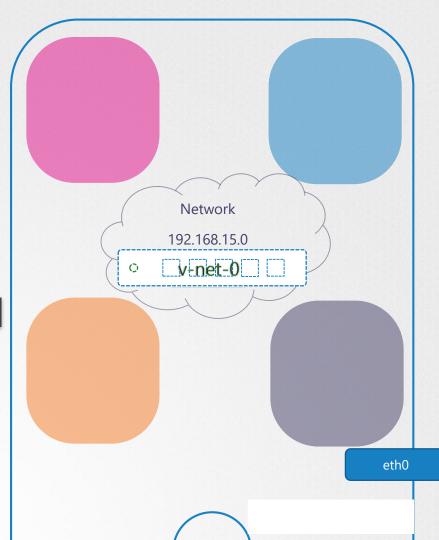


ip link add veth-red type veth peer name veth-red-br

veth-red veth-red-br

ip link add veth-blue type veth peer name veth-blue-br

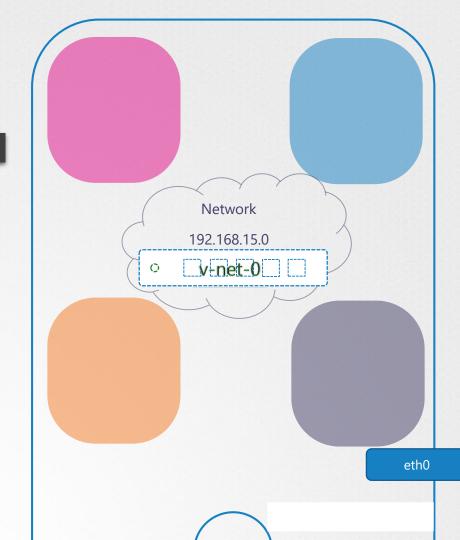
veth-blue veth-blue-br



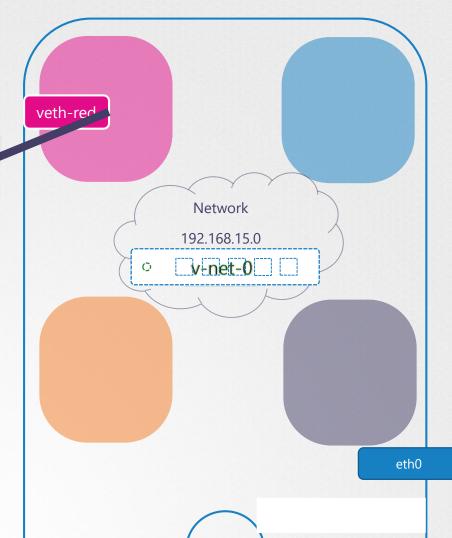
ip link set veth-red netns red

veth-red-br

veth-blue veth-blue-br



ip link set veth-red netns red ip link set veth-red-br master v-net-0 veth-red-br veth-blue veth-blue-br

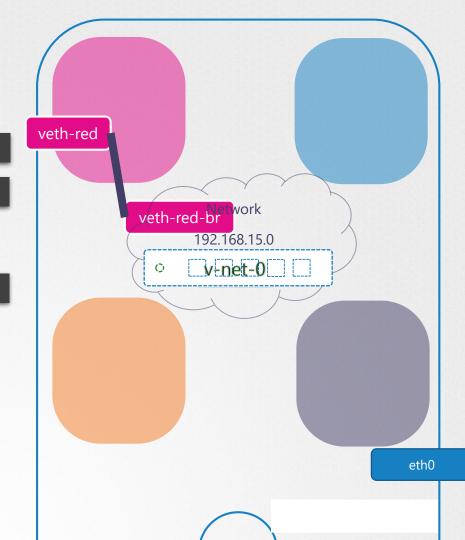


ip link set veth-red netns red

ip link set veth-red-br master v-net-0

ip link set veth-blue netns blue

veth-blue veth-blue-br



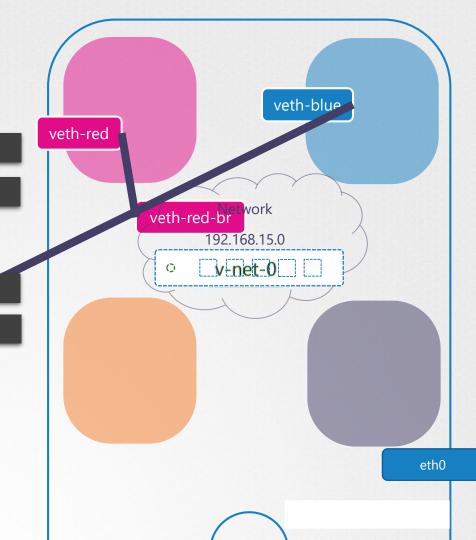
ip link set veth-red netns red

ip link set veth-red-br master v-net-0

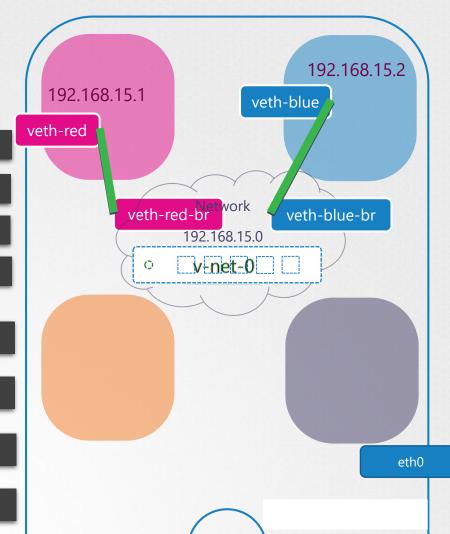
ip link set veth-blue netns blue

ip link set veth-blue-br master v-net-0

veth-blue-br



- ip link set veth-red netns red
- ip link set veth-red-br master v-net-0
- ip link set veth-blue netns blue
- ip link set veth-blue-br master v-net-0
- ip -n red addr add 192.168.15.1 dev veth-red
- ip -n blue addr add 192.168.15.2 dev veth-blue
- ip -n red link set veth-red up
- ip -n blue link set veth-blue up



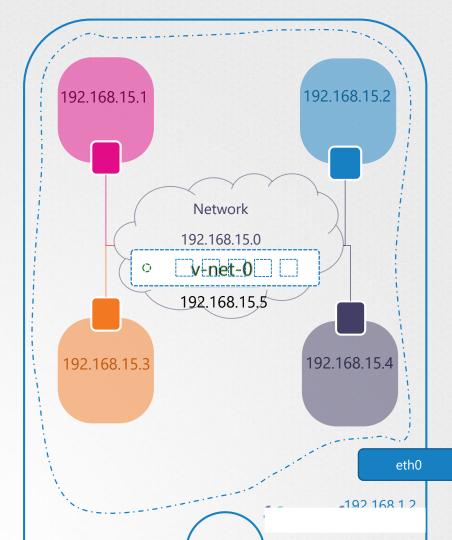
ping 192.168.15.1

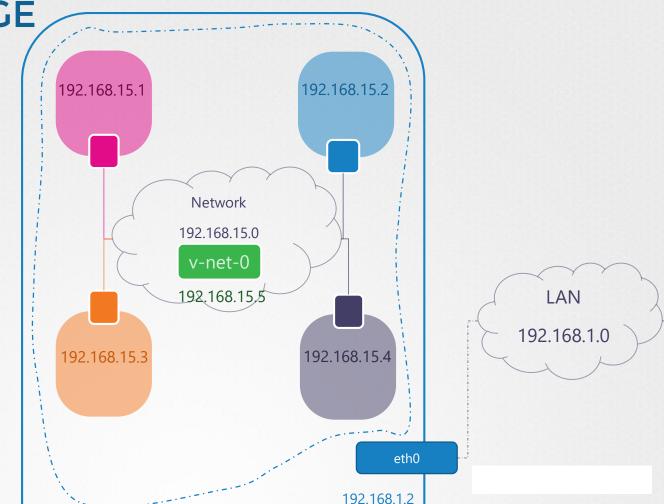
Not Reachable!

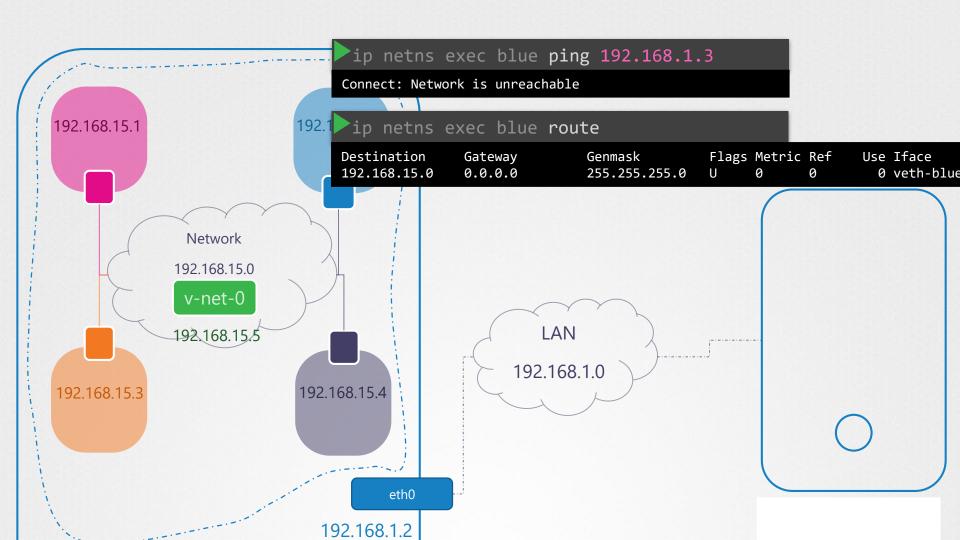
ip addr add 192.168.15.5/24 dev v-net-0

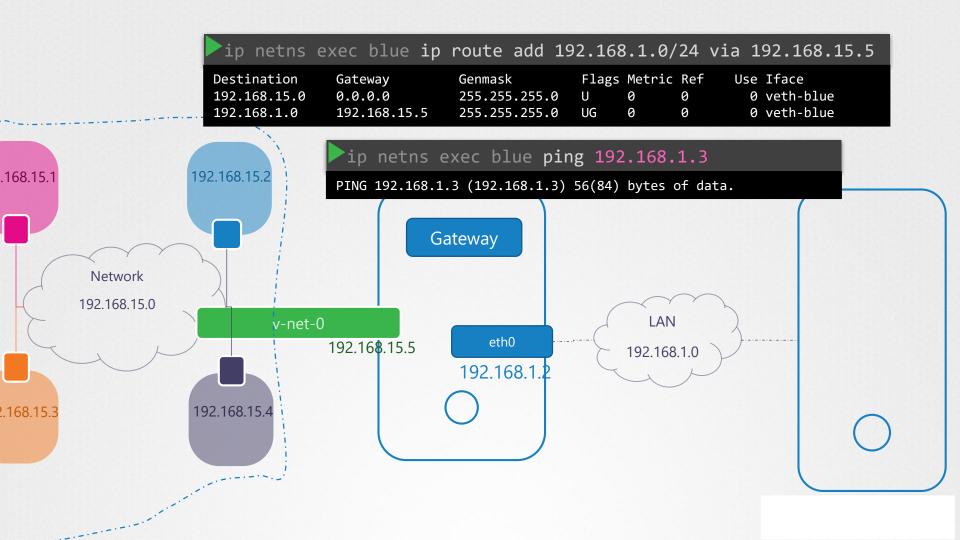
ping 192.168.15.1

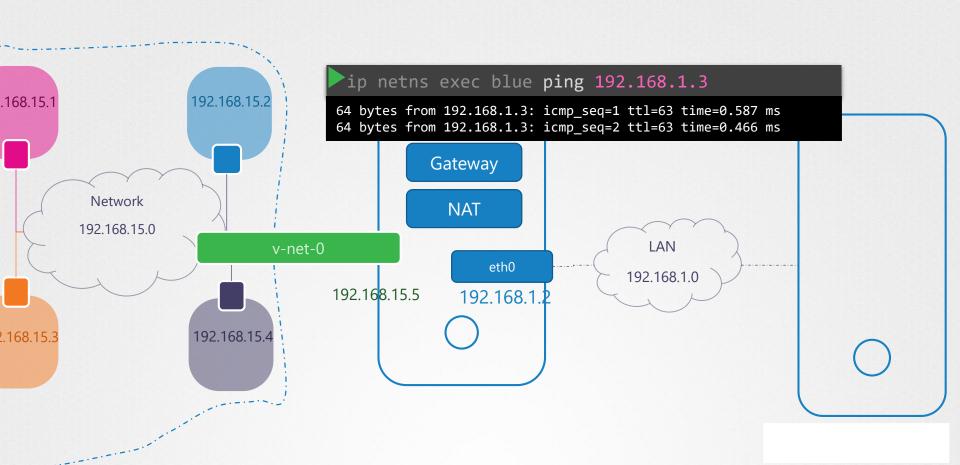
PING 192.168.15.1 (192.168.15.1) 56(84) bytes of data. 64 bytes from 192.168.15.1: icmp_seq=1 ttl=64 time=0.026 ms

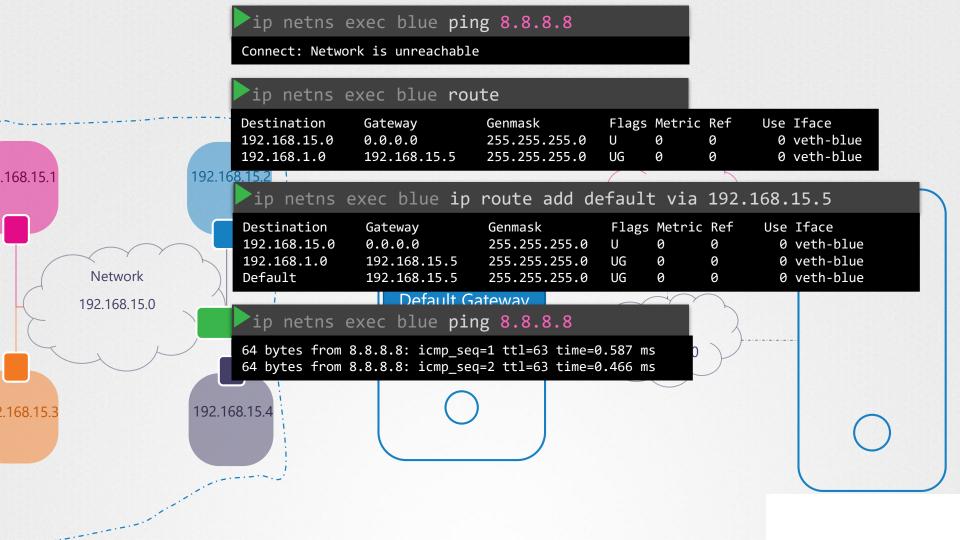


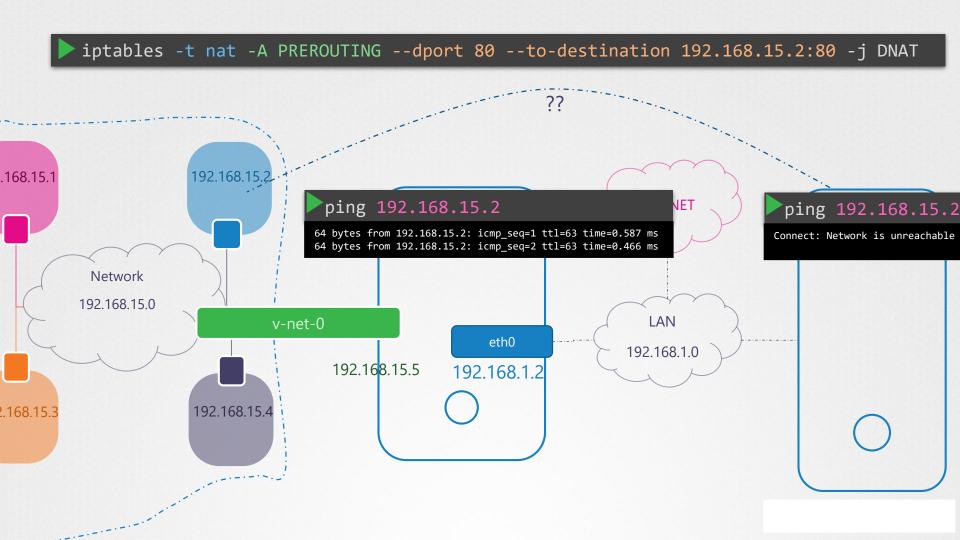












```
# Create veth pairs
ip link add veth-red type veth peer name veth-blue
# Create Add veth to respective namespaces
ip link set veth-red netns red
ip link set veth-blue netns blue
# Set IP Addresses
ip -n red addr add 192.168.1.1 dev veth-red
ip -n blue addr add 192.168.1.2 dev veth-blue
# Check IP Addresses
ip -n red addr
ip -n blue addr
# Bring up interfaces
ip -n red link set veth-red up
ip -n blue link set veth-blue up
# Bring Loopback devices up
ip -n red link set lo up
ip -n blue link set lo up
# Add default gateway
ip netns exec red ip route add default via 192.168.1.1 dev
veth-red
ip netns exec blue ip route add default via 192.168.1.2 dev
```

Create network namespaces

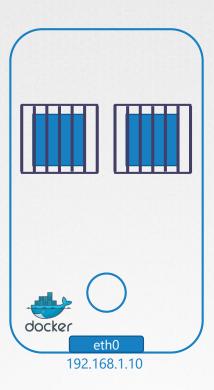
ip netns add red
ip netns add blue

```
ip netns del red
ip netns del blue
ip link del v-net-0
iptables -t nat -D POSTROUTING 1
#
ip netns add red
ip netns add blue
ip link add veth-red type veth peer name veth-red-br
ip link add veth-blue type veth peer name veth-blue-br
ip link set veth-red netns red
ip link set veth-blue netns blue
ip -n red addr add 192.168.15.2/24 dev veth-red
ip -n blue addr add 192.168.15.3/24 dev veth-blue
brctl addbr v-net-0
ip link set dev v-net-0 up
ip link set veth-red-br up
```



DOCKER NETWORKING

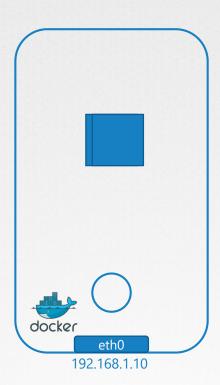
NONE



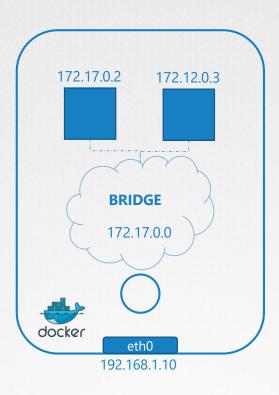
- docker run --network none nginx
- docker run --network none nginx

HOST

http://192.168.1.10:80



- docker run --network host nginx
 - docker run --network host nginx



- docker run nginx
- docker run nginx



DOCKER NETWORKING Deep Dive

docker network ls

| NETWORK ID | NAME | DRIVER | SCOPE |
|--------------|--------|--------|-------|
| 2b60087261b2 | bridge | bridge | local |
| 0beb4870b093 | host | host | local |
| 99035e02694f | none | null | local |

ip link

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
 link/loopback 00:00:00:00:00 brd 00:00:00:00:00:00
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
 link/ether 02:42:ac:11:00:08 brd ff:ff:ff:ff:
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode DEFAULT group default
 link/ether 02:42:88:56:50:83 brd ff:ff:ff:ff:ff:

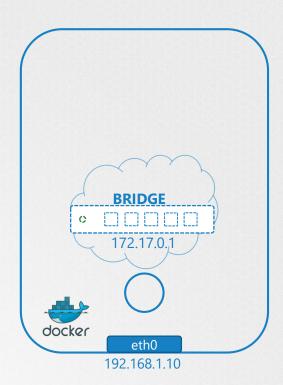
ip link add docker0 type bridge **BRIDGE** 172.17.0.1 docker eth0 192.168.1.10

ip link

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc
fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 02:42:ac:11:00:08 brd ff:ff:ff:ff:
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc
noqueue state DOWN mode DEFAULT group default
    link/ether 02:42:88:56:50:83 brd ff:ff:ff:ff:ff
```

ip addr

```
3: docker0: <NO-CARRIER, BROADCAST, MULTICAST, UP> mtu 1500 qdisc
noqueue state DOWN group default
link/ether 02:42:88:56:50:83 brd ff:ff:ff:ff:ff
inet 172.17.0.1/24 brd 172.17.0.255 scope global docker0
valid_lft forever preferred_lft forever
```



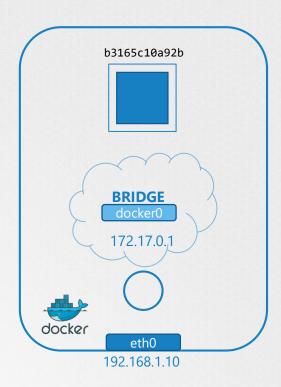
ip addr

```
3: docker0: <NO-CARRIER, BROADCAST, MULTICAST, UP> mtu 1500 qdisc
noqueue state DOWN group default
link/ether 02:42:88:56:50:83 brd ff:ff:ff:ff:ff
inet 172.17.0.1/24 brd 172.17.0.255 scope global docker0
valid_lft forever preferred_lft forever
```

ip netns

b3165c10a92b

docker inspect 942d70e585b2



docker run nginx

2e41deb9ef1b8b3d141c7bb55d883541b4

ip netns

b3165c10a92b

vethbb1c343@if7

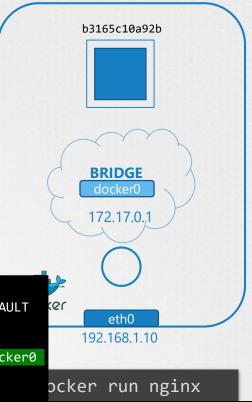
ip link

4: docker0: <BROADCAST, MULTICAST, UP, LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT group default

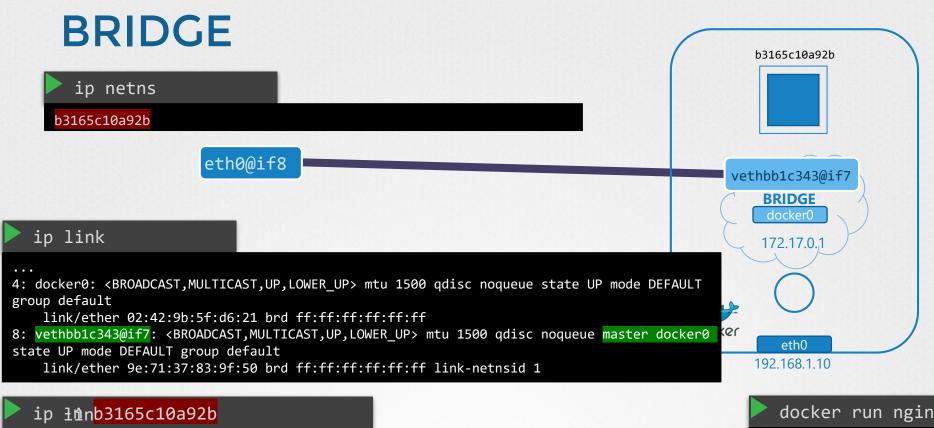
link/ether 02:42:9b:5f:d6:21 brd ff:ff:ff:ff:ff

8: vethbb1c343@if7: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master docker0 state UP mode DEFAULT group default

link/ether 9e:71:37:83:9f:50 brd ff:ff:ff:ff:ff:ff link-netnsid 1



2e41deb9ef1b8b3d141c7bb55d883541b4



7: etho@if8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT group default

link/ether 02:42:ac:11:00:03 brd ff:ff:ff:ff:ff link-netnsid 0

docker run nginx

2e41deb9ef1b8b3d141c7bb55

ip netns

b3165c10a92b

```
ip link
```

4: docker0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT group default

link/ether 02:42:9b:5f:d6:21 brd ff:ff:ff:ff:ff

8: vethbb1c343@if7: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master docker0

state UP mode DEFAULT group default
 link/ether 9e:71:37:83:9f:50 brd ff:ff:ff:ff:ff link-netnsid 1

ip -n b3165c10a92b link

7: eth0@if8: k87 k87 mode DEFAULT group default

link/ether 02:42:ac:11:00:03 brd ff:ff:ff:ff:ff:ff link-netnsid 0

b3165c10a92b eth0@if8 vethbb1c343@if7 **BRIDGE** 172.17.0.1 eth0 192.168.1.10

docker run nginx

ip netns

b3165c10a92b

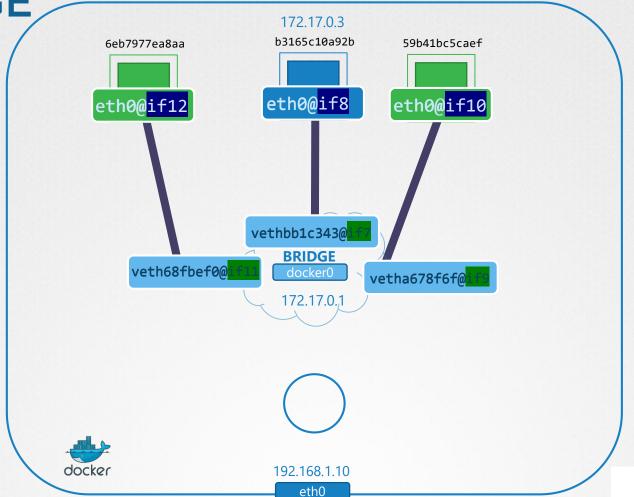
ip anddbr3165c10a92b

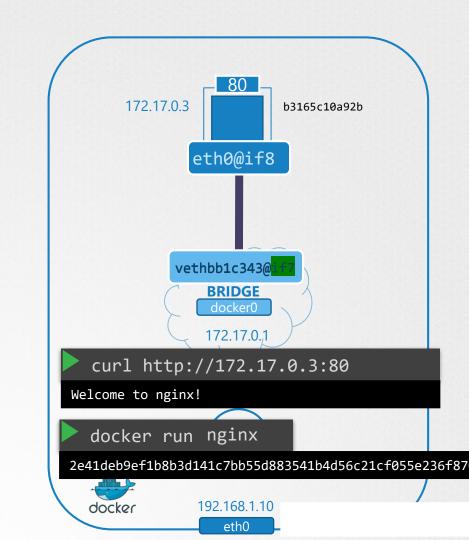
7: eth0@if8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default link/ether 02:42:ac:11:00:03 brd ff:ff:ff:ff:ff link-netnsid 0

inet 172.17.0.3/16 brd 172.17.255.255 scope global eth0
 valid_lft forever preferred_lft forever

b3165c10a92b 172.17.0.3 eth0@if8 vethbb1c343@if7 **BRIDGE** 172.17.0.1 docker eth0 192.168.1.10

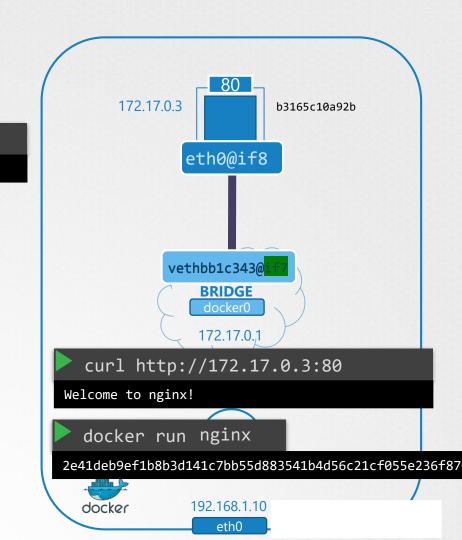
> docker run nginx 2e41deb9ef1b8b3d141c7bb55





curl http://172.17.0.3:80

curl: (7) Failed to connect... No route to host



curl http://172.17.0.3:80

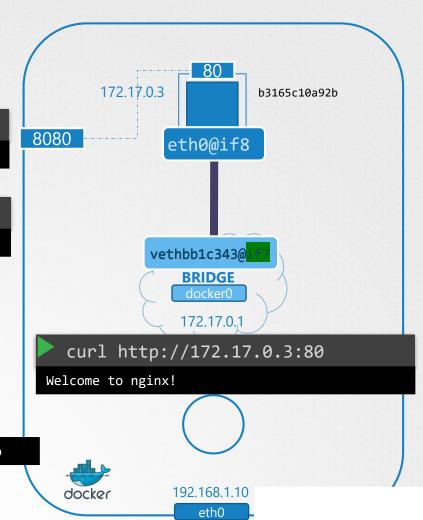
curl: (7) Failed to connect... No route to host

curl http://192.168.1.10:8080

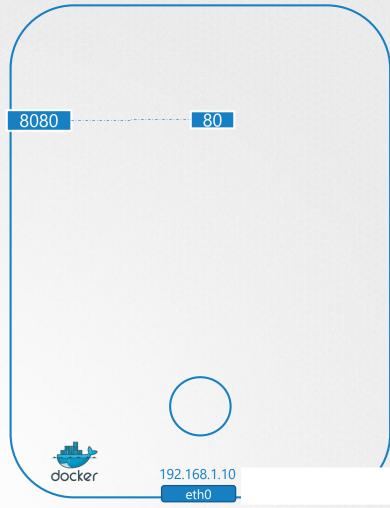
Welcome to nginx!

docker run 🗝 📆 📆 🛠 80:80

2e41deb9ef1b8b3d141c7bb55d883541b4d56c21cf055e236f870bd0f274e52b

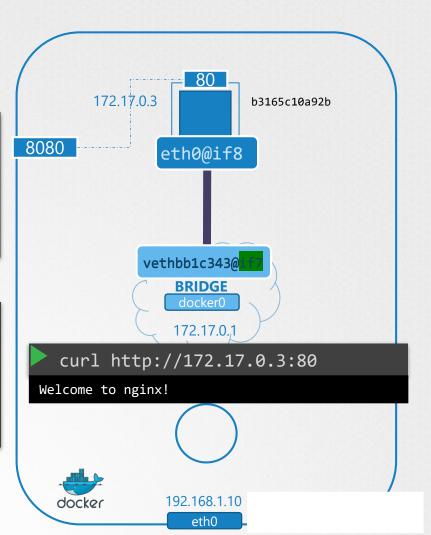


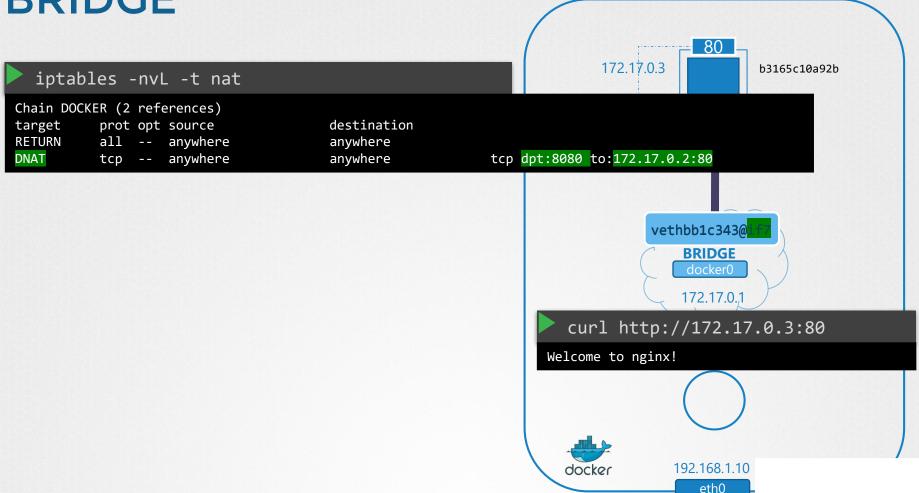
```
iptables \
    -t nat \
    -A PREROUTING \
    -j DNAT \
    --dport 8080 \
    --to-destination 80
```



```
iptables \
-t nat \
-A PREROUTING \
-j DNAT \
--dport 8080 \
--to-destination 80
```

```
iptables \
    -t nat \
    -A DOCKER \
    -j DNAT \
    --dport 8080 \
    --to-destination 172.17.0.3:80
```





https://docs.docker.com/v17.09/engine/userguide/networking/default_network/container-communication/#communicating-to-the-outside-world

In -s /var/run/docker/netns /var/run

sudo route add 172.17.0.6/32 gateway 192.168.176.14 enp0s8

sudo ip neighbor add 172.17.0.6 lladdr 02:42:ac:11:00:06 dev enp0s8

sudo bridge fdb add 02:42:ac:11:00:06 dev enp0s8 self

sudo iptables -P FORWARD ACCEPT



Container Networking Interface (CNI)

Network Namespaces

1 Create Natural Namespase

2. Create Bridge Network/Interface

Create VETH Pairs (Pine, Virtual Cable)

4 Attach vEth to Namornaco

5. Attach Other vEth to Bridge

6 Assign IP Addresses

7. Bring the interfaces up

8 Enable NAT – IP Masquerade



Create Network Namespace

2 Court Bilder Network flotterform

3. Create VETH Pairs (Pipe, Virtual Cable

4. Attach vEth to Namespace

5. Attach Other vEth to Bridge

Assign IP Addresses

7 Bring the interfaces un

Enable NAT - IP Macquerad



1. Create Network Namespa

2 Create Bridge Network/Interface

3. Create VETH Pairs (Pipe, Virtual Cable

4 Attach vEth to Namespace

5 Attach Other vEth to Bridge

6 Accion ID Addresses

7 Bring the interfaces un

8. Enable NAT – IP Masquerade



Create Network Namespace

2. Create Bridge Network/Interface

2 Create VETH Bairs (Dine Virtual Cale

4 Attach vEth to Namesna

5. Attach Other vEth to Bridge

6 Δesian IP Δda

7. Bring the interfaces up

8. Enable NAT – IP Masquerade



Create Network Namespace

2. Create Bridge Network/Interface

3 Create VETH Pairs (Pine, Virtual Cable)

4. Attach vEth to Namesnace

Attach Other vEth to Bridge

6 Assign IP Addres

7. Bring the interfaces u

8 Enable NAT - IP Masquerade









Create Network Namespace

1. Create Network Namespace

1. Create Network Namespace

1. Create Network Namespace

1. Create Network Namespace

bridge add <cid> <namespace>

bridge add <cid> <namespace>

bridge add 2e34dcf34 /var/run/netn<u>s/2e34dcf34</u>

- 2. Create Bridge Network/Interface
- 3. Create VETH Pairs (Pipe, Virtual Cable)
- 4. Attach vEth to Namespace
- 5. Attach Other vEth to Bridge
- 6. Assign IP Addresses
- 7. Bring the interfaces up
- 8. Enable NAT IP Masquerade









Create Network Namespace

1. Create Network Namespace

1. Create Network Namespace

1. Create Network Namespace



CONTAINER NETWORK INTERFACE

BRIDGE 2. Create Bridge Network/Interface 3. Create VETH Pairs (Pipe, Virtual Cable) 4. Attach vEth to Namespace 5. Attach Other vEth to Bridge 6. Assign IP Addresses 7. Bring the interfaces up 8. Enable NAT – IP Masquerade



CONTAINER NETWORK INTERFACE

- Container Runtime must create network namespace
- ☐ Identify network the container must attach to
- Container Runtime to invoke Network Plugin (bridge) when container is ADDed.
- Container Runtime to invoke Network Plugin (bridge) when container is DELeted.
- ☐ JSON format of the Network Configuration





- ☐ Must support command line arguments ADD/DEL/CHECK
- Must support parameters container id, network ns etc...
- Must manage IP Address assignment to PODs
- ☐ Must Return results in a specific format





CONTAINER NETWORK INTERFACE

docker run --network=cni-bridge nginx

docker run --network=none nginx

bridge add 2e34dcf34 /var/run/netns/2e34dcf34









CONTAINER NETWORK MODEL (CNM)

BRIDGE

VLAN

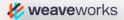
IPVLAN

MACVLAN

WINDOWS

DHCP

host-local











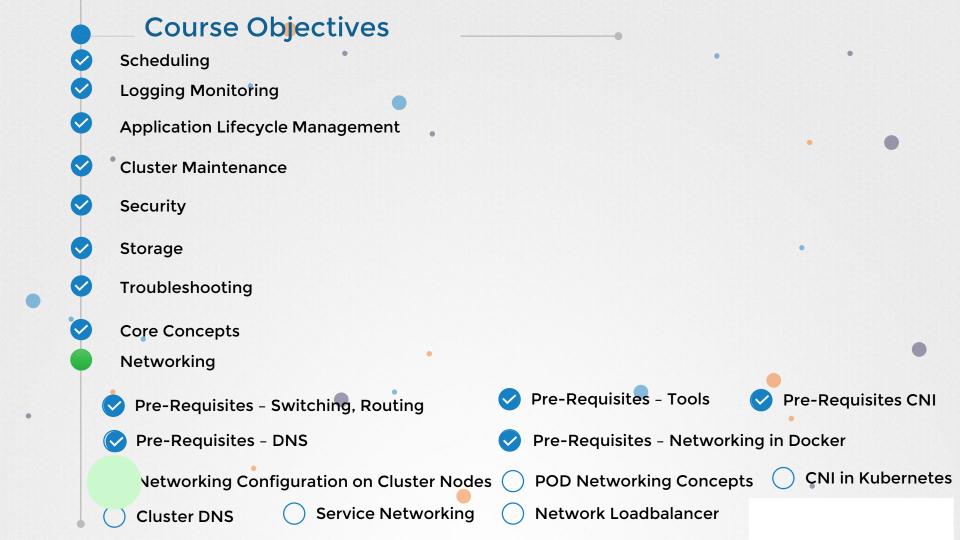






docker run --network=none nginx

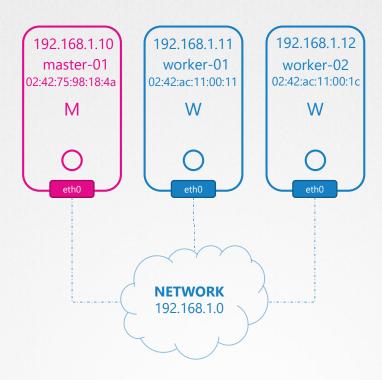
bridge add 2e34dcf34 /var/run/netns/2e34dcf34



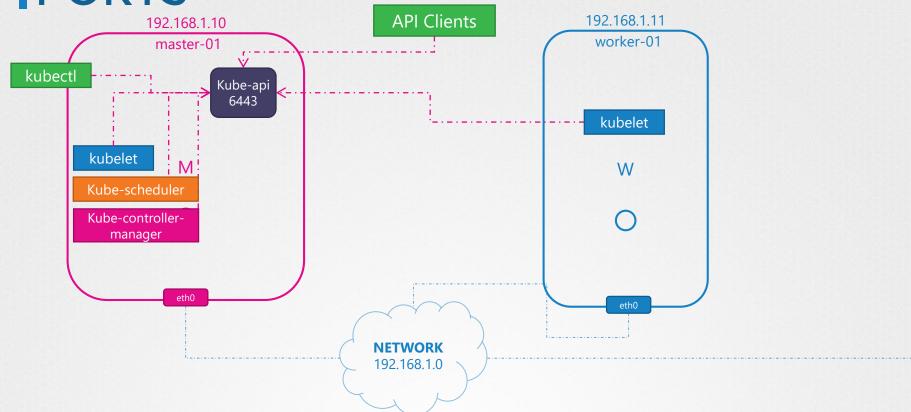


Networking Cluster Nodes

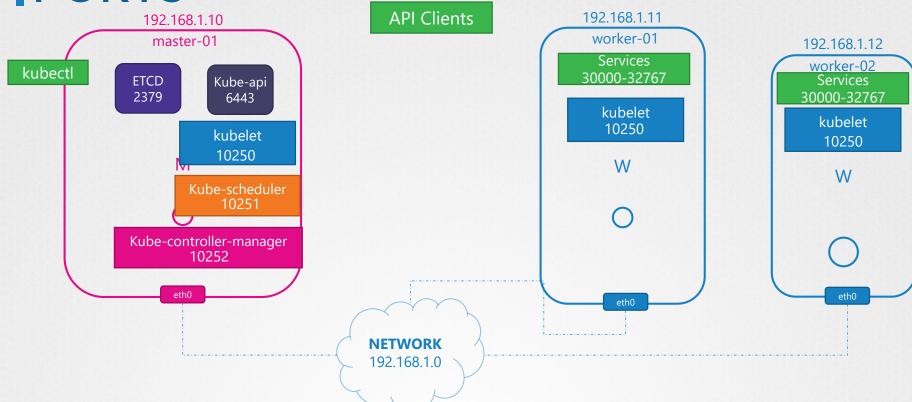
IIP & FQDN

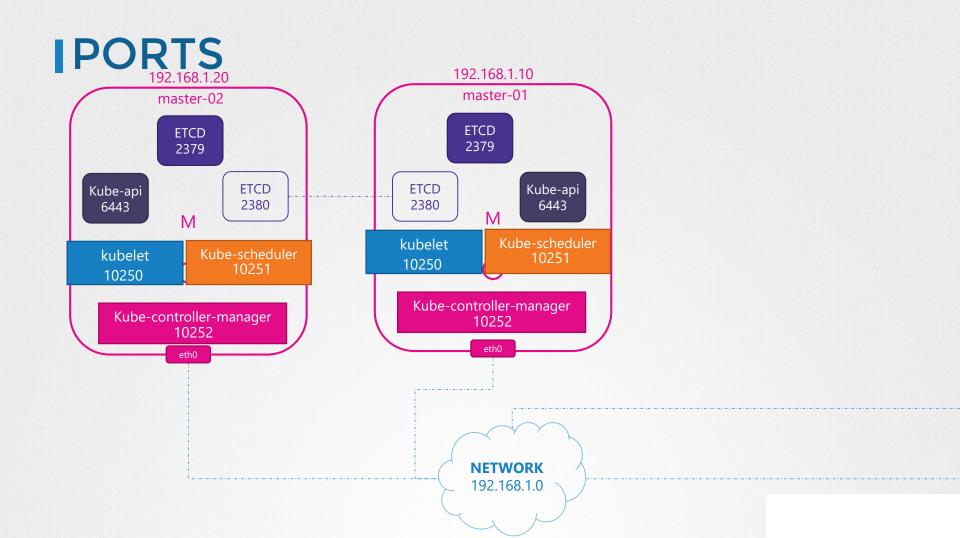


IPORTS



IPORTS





Documentation

Check required ports

Master node(s) 🔗

| Protocol | Direction | Port Range | Purpose | Used By |
|----------|-----------|------------|-------------------------|----------------------|
| TCP | Inbound | 6443* | Kubernetes API server | All |
| TCP | Inbound | 2379-2380 | etcd server client API | kube-apiserver, etcd |
| TCP | Inbound | 10250 | Kubelet API | Self, Control plane |
| TCP | Inbound | 10251 | kube-scheduler | Self |
| TCP | Inbound | 10252 | kube-controller-manager | Self |

Worker node(s)

| Protocol | Direction | Port Range | Purpose | Used By |
|----------|-----------|-------------|---------------------|---------------------|
| TCP | Inbound | 10250 | Kubelet API | Self, Control plane |
| TCP | Inbound | 30000-32767 | NodePort Services** | All |

https://kubernetes.io/docs/setup/independent/install-kubeadm/#check-required

ICOMMMANDS

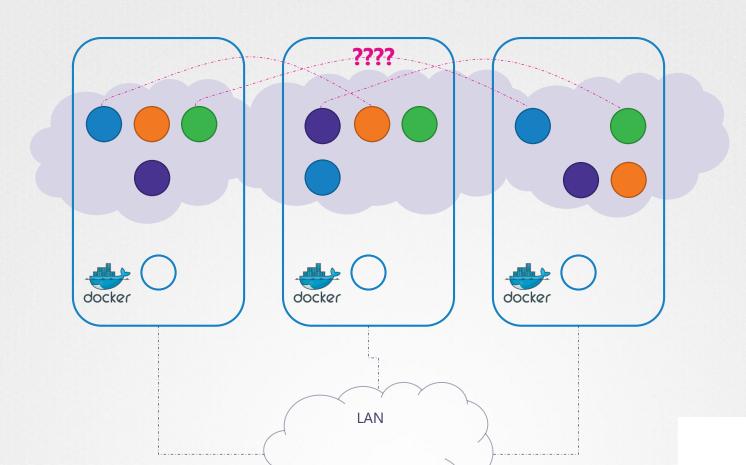
```
ip link
 ip addr
 ip addr add 192.168.1.10/24 dev eth0
 ip route
ip route add 192.168.1.0/24 via 192.168.2.1
 cat /proc/sys/net/ipv4/ip_forward
 arp
 netstat -plnt
```

route

Course Objectives Scheduling **Logging Monitoring Application Lifecycle Management Cluster Maintenance** Security Storage **Troubleshooting Core Concepts** Networking Pre-Requisites - Tools **Pre-Requisites CNI** Pre-Requisites - Switching, Routing Pre-Requisites - Networking in Docker **Pre-Requisites - DNS CNI** in Kubernetes **Networking Configuration on Cluster Node: POD Networking Concepts Service Networking Network Loadbalancer Cluster DNS**

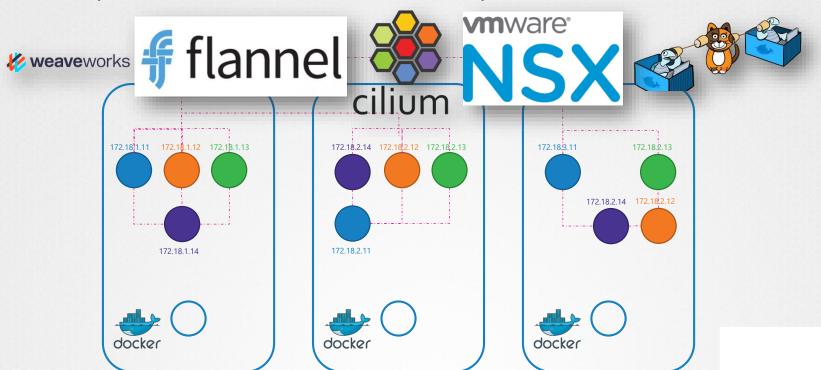


POD Networking Concepts



Networking Model

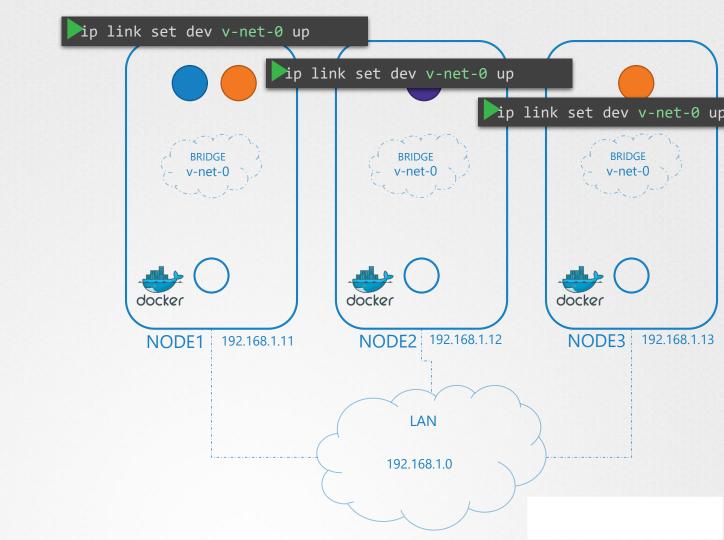
- Every POD should have an IP Address
- Every POD should be able to communicate with every other POD in the same node.
- Every POD should be able to communicate with every other POD on other nodes without NAT.

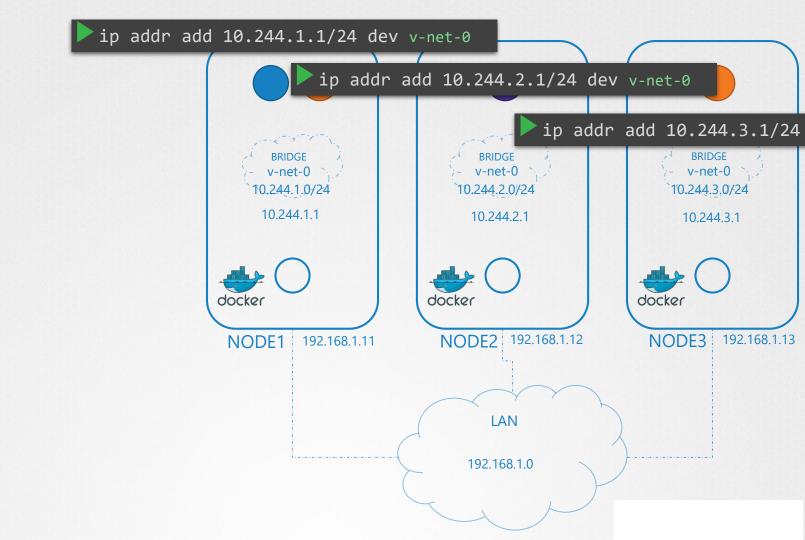


ip link add v-net-0 type bridge Vetworking Voce ip link set dev v-net-0 up

- Every POD should have an IP Address
- Every POD shou pip addr add 192.168.15.5/24 dev v-net-0 me node.
- Every POD should be able to communicate with every other POD on other nodes without NAT.

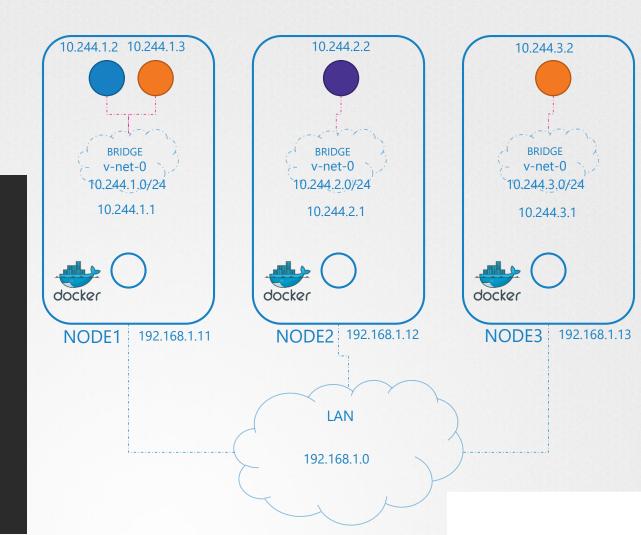
ip link add veth-red type veth peer name veth-red-br ip link set veth-red netns red pip -n red addr add 192.168.15.1 dev veth-red ip -n red link set veth-red up ip link set veth-red-br master v-net-0 ip netns exec blue ip route add 192.168.1.0/24 via 192.168.15.5 iptables -t nat -A POSTROUTING -s 192.168.15.0/24 -j MASQUERADE docker

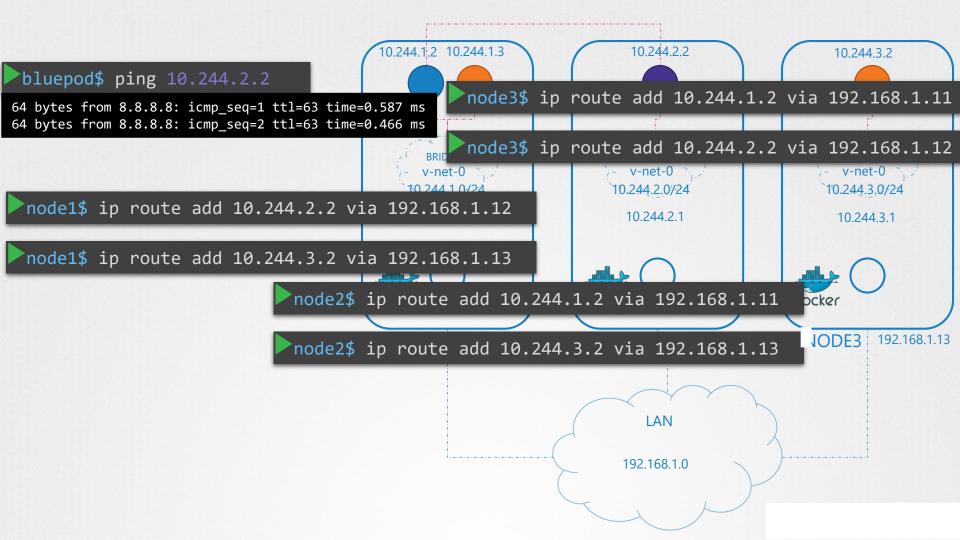


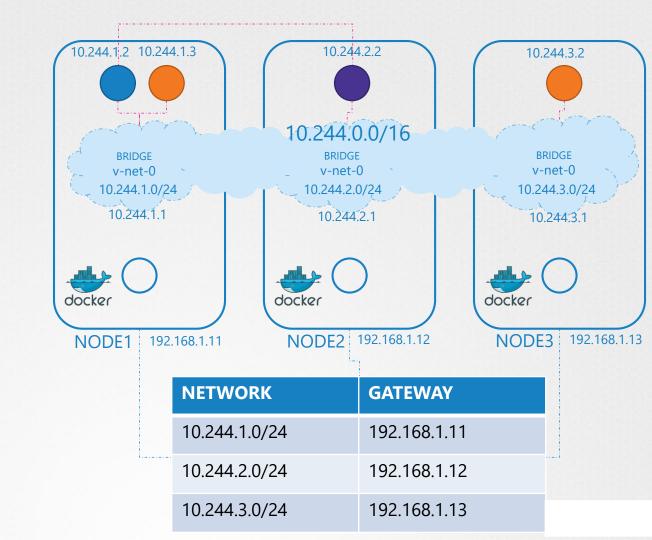


net-script.sh

```
# Create veth pair
ip link add .....
# Attach veth pair
ip link set .....
ip link set .....
# Assign IP Address
ip -n <namespace> addr add .....
ip -n <namespace> route add .....
ip -n <namespace> link set .....
```











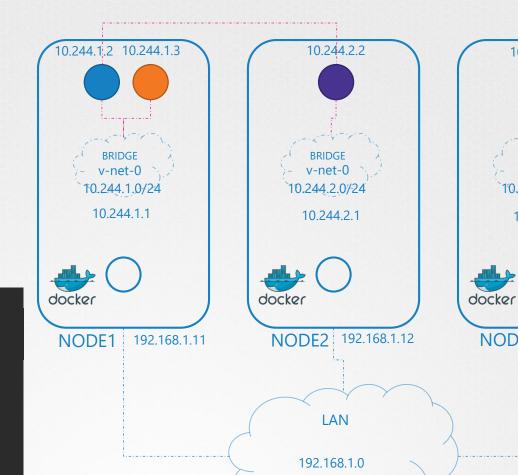
(CNI)

net-script.sh

ip link add

Attach veth pair ip link set ip link set

Assign IP Address ip -n <namespace> addr add ip -n <namespace> route add



10.244.3.2

BRIDGE

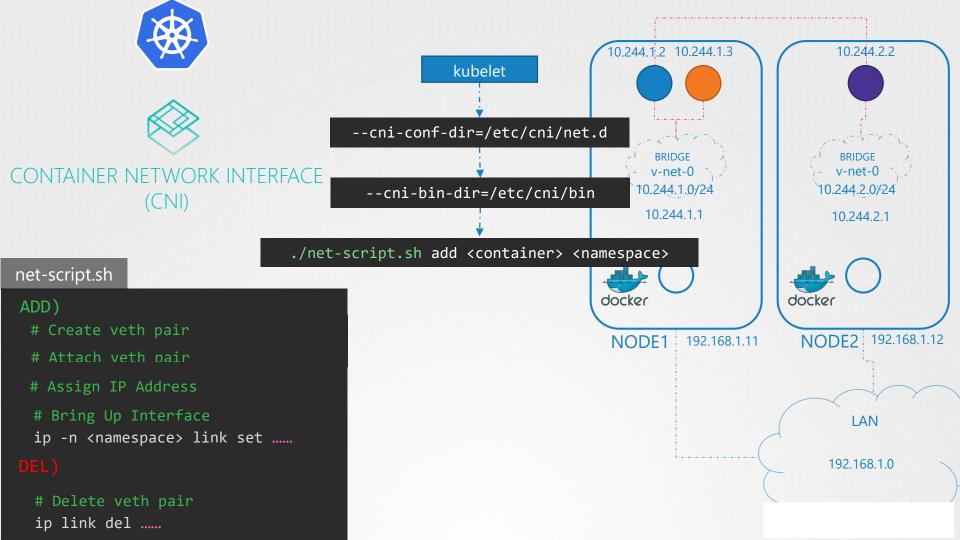
v-net-0

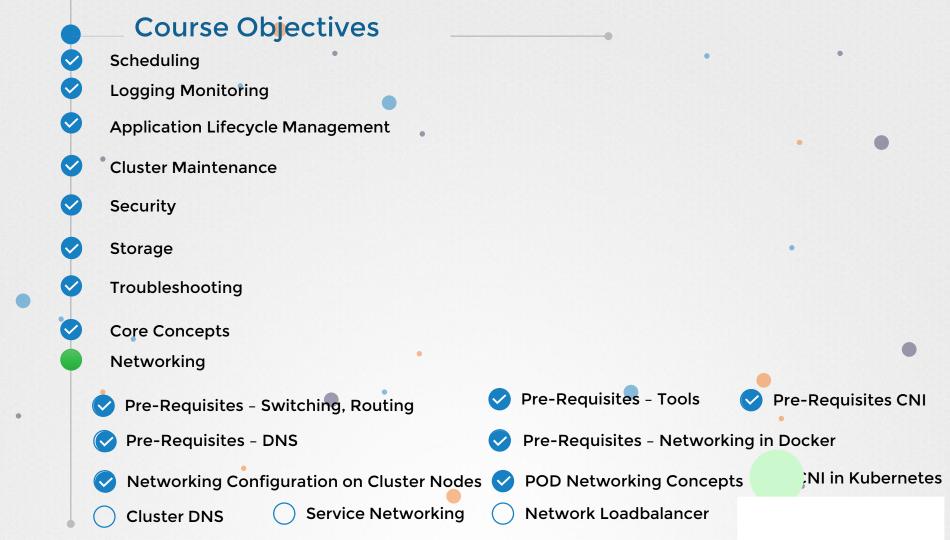
10.244.3.0/24

10.244.3.1

NODE3

192.168.1.13







Container Networking Interface (CNI) IN KUBERNETES

IPre-Requisites

- ✓ Network Namespaces in Linux
- ✓ Networking in Docker
- ✓ Why and what is Container Network Interface (CNI)?
- ✓ CNI Plugins



- ✓ Container Runtime must create network namespace
- ✓ Identify network the container must attach to
- ✓ Container Runtime to invoke Network Plugin (bridge) when container is ADDed.
- ✓ Container Runtime to invoke Network Plugin (bridge) when container is DELeted.
- ✓ JSON format of the Network Configuration



| Configuring CNI

kubelet.service

```
ExecStart=/usr/local/bin/kubelet \\
    --config=/var/lib/kubelet/kubelet-config.yaml \\
    --container-runtime=remote \\
    --container-runtime-endpoint=unix:///var/run/containerd/containerd.sock \\
    --image-pull-progress-deadline=2m \\
    --kubeconfig=/var/lib/kubelet/kubeconfig \\
    --network-plugin=cni \\
    --cni-bin-dir=/opt/cni/bin \\
    --cni-conf-dir=/etc/cni/net.d \\
    --register-node=true \\
    --v=2
```

| View kubelet options

```
ps -aux | grep kubelet

root 2095 1.8 2.4 960676 98788 ? Ssl 02:32 0:36 /usr/bin/kubelet --bootstrap-
kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --
config=/var/lib/kubelet/config.yaml --cgroup-driver=cgroupfs --cni-bin-dir=/opt/cni/bin --cni-
conf-dir=/etc/cni/net.d --network-plugin=cni
```

ls /opt/cni/bin

bridge dhcp flannel host-local ipvlan loopback macvlan portmap ptp sample tuning vlan weave-ipam weave-net weave-plugin-2.2.1

ls /etc/cni/net.d

10-bridge.conf

View kubelet options

ls /etc/cni/net.d

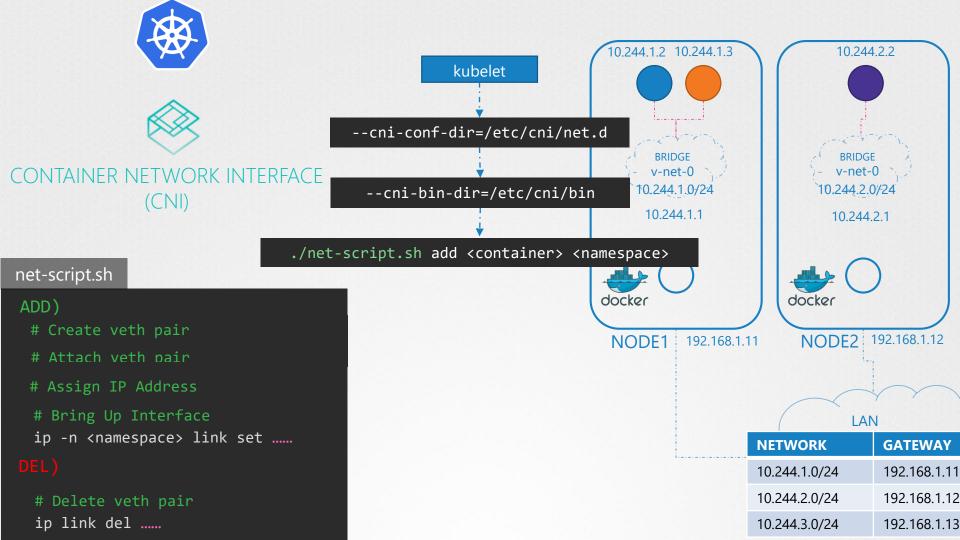
10-bridge.conf

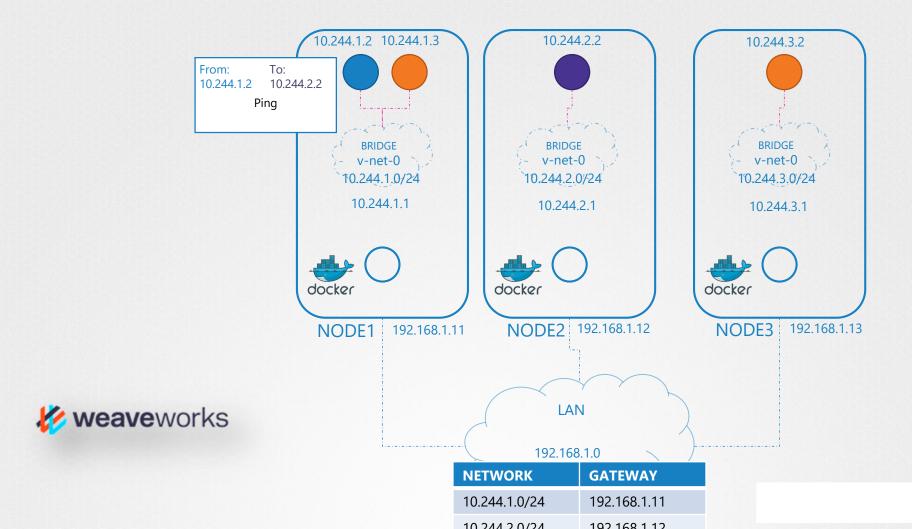
```
cat /etc/cni/net.d/10-bridge.conf
"isGateway": true,
"ipMasq": true,
```

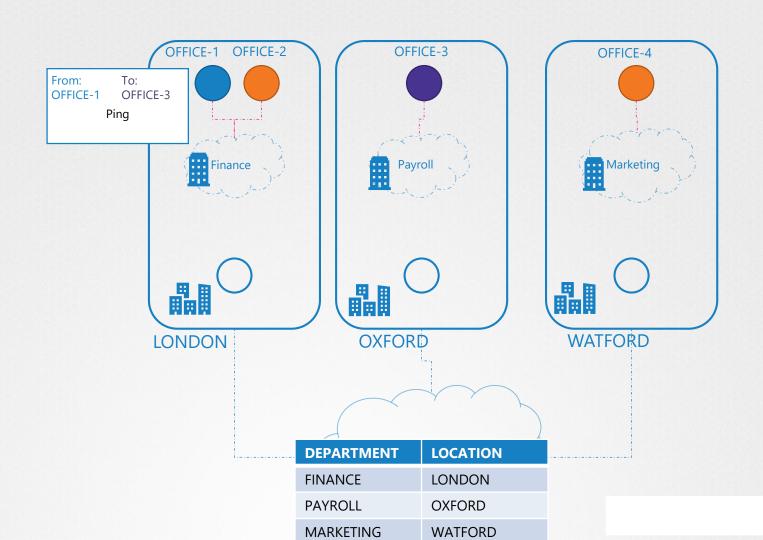
Course Objectives Scheduling **Logging Monitoring Application Lifecycle Management Cluster Maintenance** Security Storage **Troubleshooting Core Concepts** Networking **Pre-Requisites - Tools Pre-Requisites CNI** Pre-Requisites - Switching, Routing Pre-Requisites - DNS, IPAM, Firewalls, LBs Pre-Requisites - Networking in Docker **CNI** in Kubernetes **Networking Configuration on Cluster Nodes** (**POD Networking Concepts Service Networking Network Loadbalancer Cluster DNS**

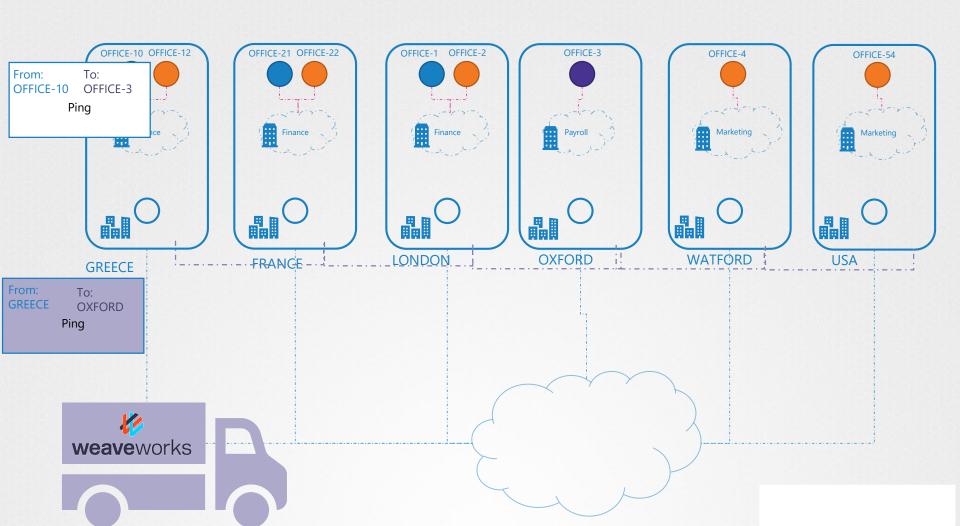


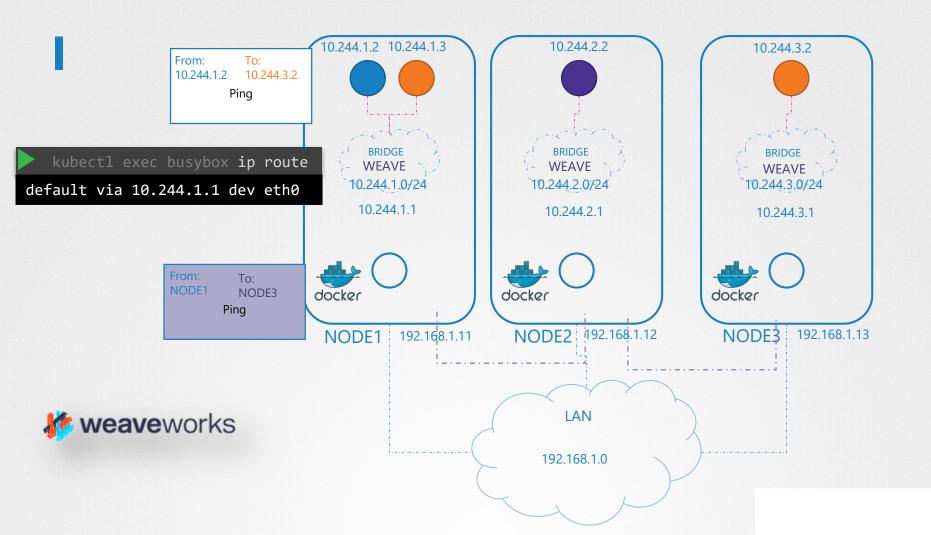
WeaveWorks (CNI)











IDeploy Weave

```
kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')"

serviceaccount/weave-net created clusterrole.rbac.authorization.k8s.io/weave-net created clusterrolebinding.rbac.authorization.k8s.io/weave-net created role.rbac.authorization.k8s.io/weave-net created rolebinding.rbac.authorization.k8s.io/weave-net created daemonset.extensions/weave-net created
```

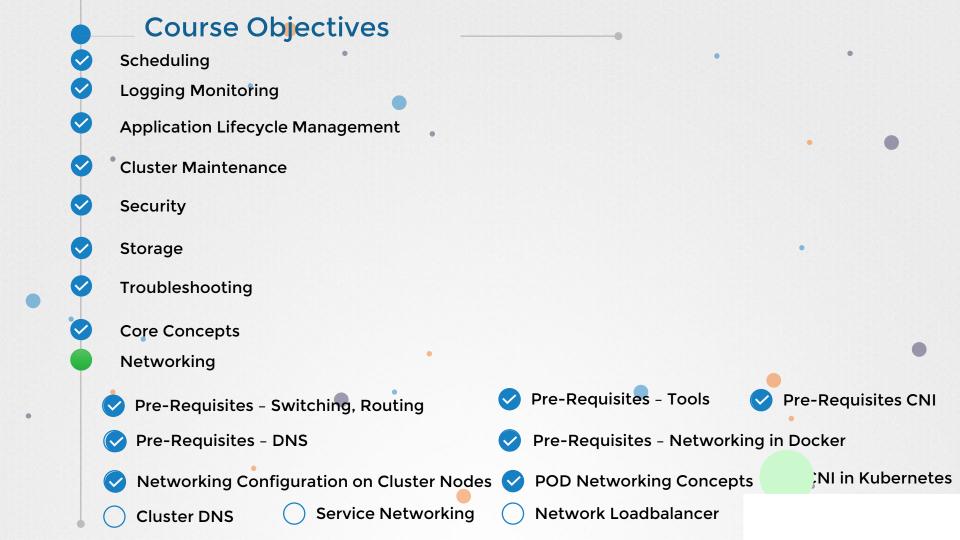
Weave Peers

| kubectl | σ _Δ + | node | _n | kuha-sv | /ctam |
|---------|------------------|------|-----|---------|-------------|
| KUDECLI | get | pous | -11 | Kube-Sy | / S L E III |

| NAME | READY | STATUS | RESTARTS | AGE | IP | NODE | |
|--------------------------------------------|-------|---------|----------|-------------|-------------|--------|---------------|
| NOMINATED NODE coredns-78fcdf6894-99khw | 1/1 | Running | 0 | 1 9m | 10.44.0.2 | master | <none></none> |
| coredns-78fcdf6894-p7dpj | 1/1 | Running | 0 | 19m | 10.44.0.1 | master | <none></none> |
| etcd-master | 1/1 | Running | 0 | 18m | 172.17.0.11 | master | <none></none> |
| kube-apiserver-master | 1/1 | Running | 0 | 18m | 172.17.0.11 | master | <none></none> |
| kube-scheduler-master | 1/1 | Running | 0 | 17m | 172.17.0.11 | master | <none></none> |
| weave-net-5gcmb | 2/2 | Running | 1 | 19m | 172.17.0.30 | node02 | <none></none> |
| weave-net-fr9n9 | 2/2 | Running | 1 | 19m | 172.17.0.11 | master | <none></none> |
| weave-net-mc6s2 | 2/2 | Running | 1 | 19m | 172.17.0.23 | node01 | <none></none> |
| weave-net-tbzvz | 2/2 | Running | 1 | 19m | 172.17.0.52 | node03 | <none></none> |

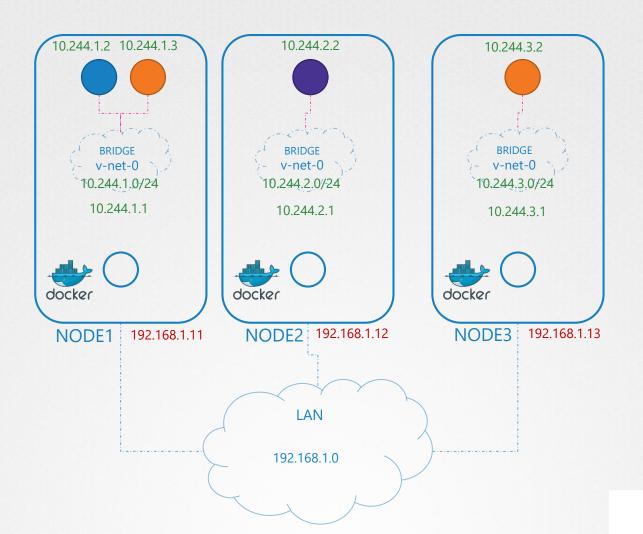
kubectl logs weave-net-5gcmb weave -n kube-system

```
INFO: 2019/03/03 03:41:08.643858 Command line options: map[status-addr:0.0.0.0:6782 http-addr:127.0.0.1:6784 ipalloc-range:10.32.0.0/12 name:9e:96:c8:09:bf:c4 nickname:node02 conn-limit:30
datapath:datapath db-prefix:/weavedb/weave-net host-root:/host port:6783 docker-api: expect-npc:true ipalloc-init:consensus=4 no-dns:true]
INFO: 2019/03/03 03:41:08.643980 weave 2.2.1
INFO: 2019/03/03 03:41:08.751508 Bridge type is bridged fastdp
INFO: 2019/03/03 03:41:08.751526 Communication between peers is unencrypted.
INFO: 2019/03/03 03:41:08.753583 Our name is 9e:96:c8:09:bf:c4(node02)
INFO: 2019/03/03 03:41:08.753615 Launch detected - using supplied peer list: [172.17.0.11 172.17.0.23 172.17.0.30 172.17.0.52]
INFO: 2019/03/03 03:41:08.753632 Checking for pre-existing addresses on weave bridge
INFO: 2019/03/03 03:41:08.756183 [allocator 9e:96:c8:09:bf:c4] No valid persisted data
INFO: 2019/03/03 03:41:08.761033 [allocator 9e:96:c8:09:bf:c4] Initialising via deferred consensus
INFO: 2019/03/03 03:41:08.761091 Sniffing traffic on datapath (via ODP)
INFO: 2019/03/03 03:41:08.761659 ->[172.17.0.23:6783] attempting connection
INFO: 2019/03/03 03:41:08.817477 overlay switch ->[8a:31:f6:b1:38:3f(node03)] using fastdp
INFO: 2019/03/03 03:41:08.819493 sleeve ->[172.17.0.52:6783|8a:31:f6:b1:38:3f(node03)]: Effective MTU verified at 1438
INFO: 2019/03/03 03:41:09.107287 Weave version 2.5.1 is available; please update at https://github.com/weaveworks/weave/releases/download/v2.5.1/weave
INFO: 2019/03/03 03:41:09.284907 Discovered remote MAC 8a:dd:b5:14:8f:a3 at 8a:dd:b5:14:8f:a3(node01)
INFO: 2019/03/03 03:41:09.331952 Discovered remote MAC 8a:31:f6:b1:38:3f at 8a:31:f6:b1:38:3f(node03)
INFO: 2019/03/03 03:41:09.355976 Discovered remote MAC 8a:a5:9c:d2:86:1f at 8a:31:f6:b1:38:3f(node03)
```

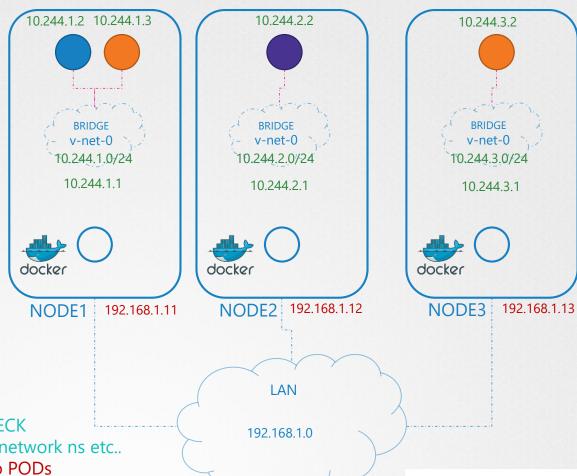




IPAM (CNI)







CNI Plugin Responsibilities:

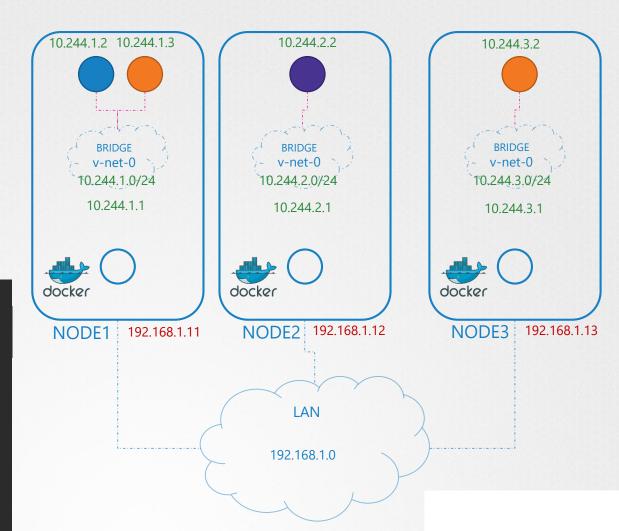
- ✓ Must support arguments ADD/DEL/CHECK
- ✓ Must support parameters container id, network ns etc..
- ✓ Must manage IP Address assignment to PODs
- ✓ Must Return results in a specific format



CONTAINER NETWORK INTERFACE (CNI)

net-script.sh

ADD) # Create veth pair # Attach veth pair # Assign IP Address # Bring Up Interface ip -n <namespace> link set DEL) # Delete veth pair ip link del

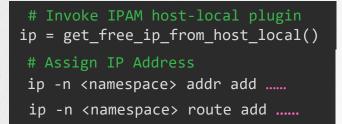


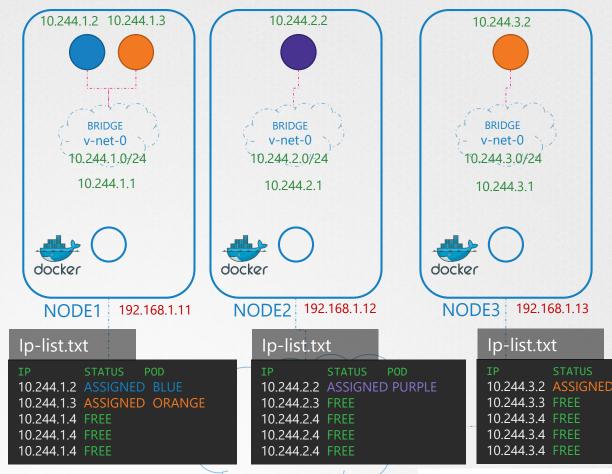


CONTAINER NETWORK INTERFACE (CNI)

DHCP

host-local



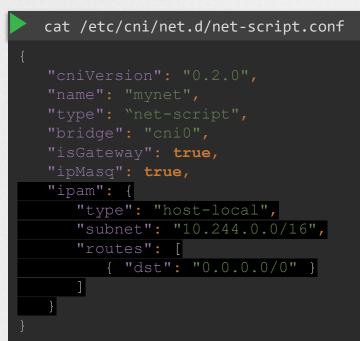


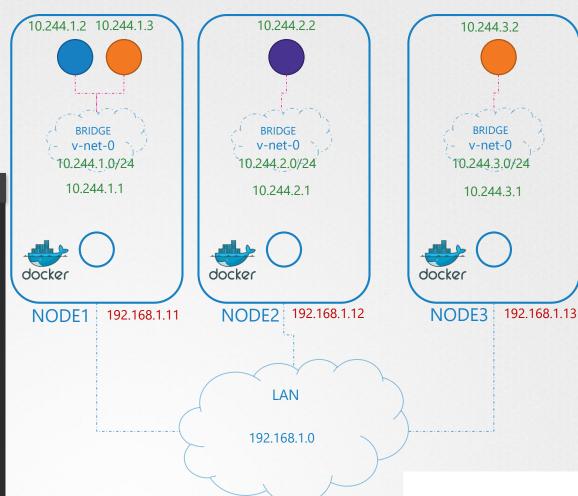


CONTAINER NETWORK INTERFACE (CNI)

DHCP

host-local



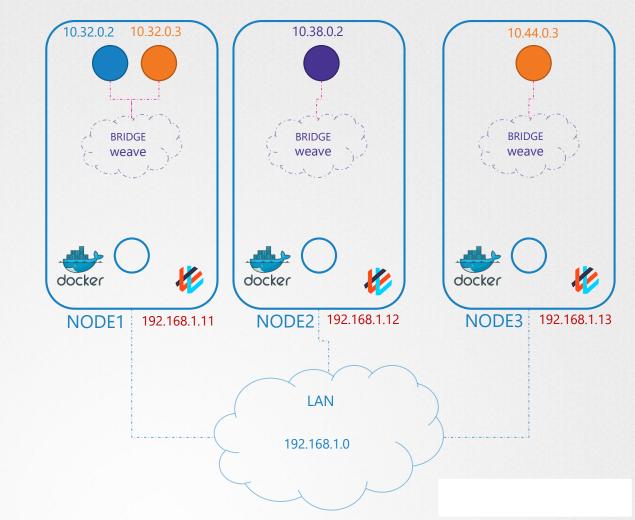


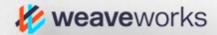
10.32.0.0/12

10.32.0.1 > 10.47.255.254

1,048,574 !!!!!

10.32.0.1 10.38.0.0 10.44.0.0





Course Objectives Scheduling **Logging Monitoring Application Lifecycle Management Cluster Maintenance** Security

- Storage **Troubleshooting**

Core Concepts

- Networking

 - Pre-Requisites Switching, Routing
 - Pre-Requisites DNS
 - Networking Configuration on Cluster Nodes <
 - **Cluster DNS Service Networking**
- **Network Loadbalancer**

Pre-Requisites - Tools

POD Networking Concepts

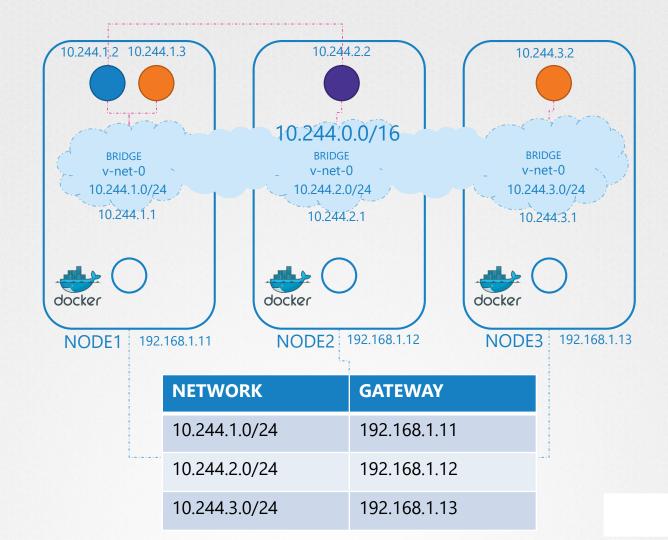
Pre-Requisites - Networking in Docker

Pre-Requisites CNI

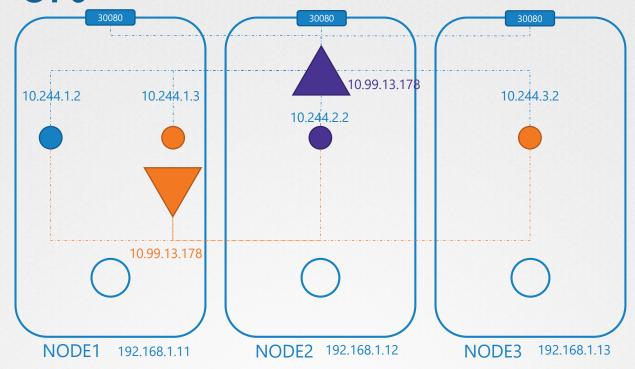
CNI in Kubernetes

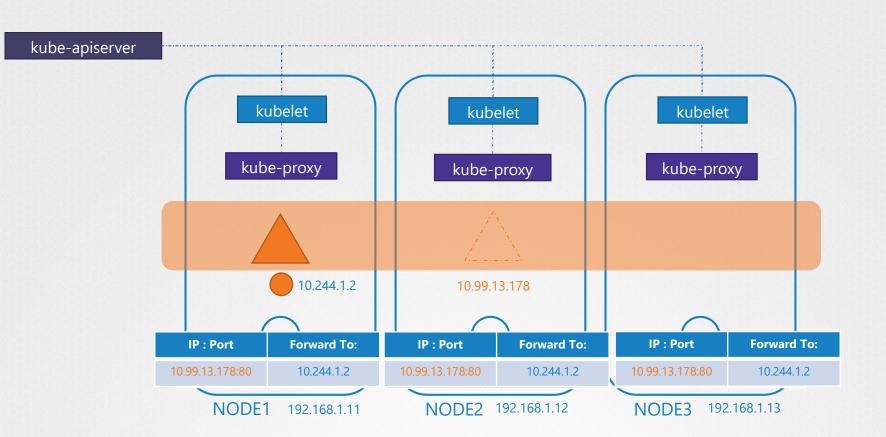


Service Networking



NodePort





userspace

iptables

ipvs

kube-proxy --proxy-mode [userspace | iptablkes | ipvs] ...

| IP : Port | Forward To: |
|-----------------|-------------|
| 10.99.13.178:80 | 10.244.1.2 |

iptables

| IP : Port | Forward To: |
|-----------------|-------------|
| 10.99.13.178:80 | 10.244.1.2 |

kubelet get pods -o wide

NAME READY STATUS RESTARTS AGE IP NODE db 1/1 Running 0 14h 10.244.1.2 node-1

kubelet get service

NAME TYPE CLUSTER-IP PORT(S) AGE db-service ClusterIP 10.103.132.104 3306/TCP 12h

10.244.1.2

10.244.0.0/16

10.244.0.0 => 10.244.255.255

10.96.0.0 => 10.111.255.255

kube-api-server --service-cluster-ip-range ipNet (Default: 10.0.0.0/24)

ps aux | grep kube-api-server

kube-apiserver --authorization-mode=Node,RBAC --service-cluster-iprange=10.96.0.0/12

iptables

| IP : Port | Forward To: |
|-----------------|-------------|
| 10.99.13.178:80 | 10.244.1.2 |

kubelet get pods -o wide

| NAME | READY | STATUS | RESTARTS | AGE | IP | NODE |
|------|-------|---------|----------|-----|------------|--------|
| db | 1/1 | Running | 0 | 14h | 10.244.1.2 | node-1 |

kubelet get service

NAME TYPE CLUSTER-IP PORT(S) AGE db-service ClusterIP 10.103.132.104 3306/TCP 12h

iptables -L -t net | grep db-service

```
KUBE-SVC-XA50GUC7YRHOS3PU tcp -- anywhere 10.103.132.104 /* default/db-service: cluster IP */ tcp dpt:3306

DNAT tcp -- anywhere anywhere /* default/db-service: */ tcp to:10.244.1.2:3306

KUBE-SEP-JBWCWHHQM57V2WN7 all -- anywhere anywhere /* default/db-service: */
```

iptables

| IP : Port | Forward To: |
|-----------------|-------------|
| 10.99.13.178:80 | 10.244.1.2 |

iptables -L -t net | grep db-service

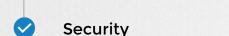
```
KUBE-SVC-XA50GUC7YRHOS3PU tcp -- anywhere 10.103.132.104 /* default/db-service: cluster IP */ tcp dpt:3306

DNAT tcp -- anywhere anywhere /* default/db-service: */ tcp to:10.244.1.2:3306

KUBE-SEP-JBWCWHHQM57V2WN7 all -- anywhere anywhere /* default/db-service: */
```

cat /var/log/kube-proxy.log

Course Objectives Scheduling **Logging Monitoring Application Lifecycle Management Cluster Maintenance** Storage



- **Troubleshooting**
 - **Core Concepts**
 - Networking
 - Pre-Requisites Switching, Routing
 - **Pre-Requisites DNS**
 - Networking Configuration on Cluster Nodes <
 - **POD Networking Concepts** Service Networking Cluster DNS
 - **Network Loadbalancer**

Pre-Requisites - Networking in Docker

Pre-Requisites - Tools

Pre-Requisites CNI

CNI in Kubernetes



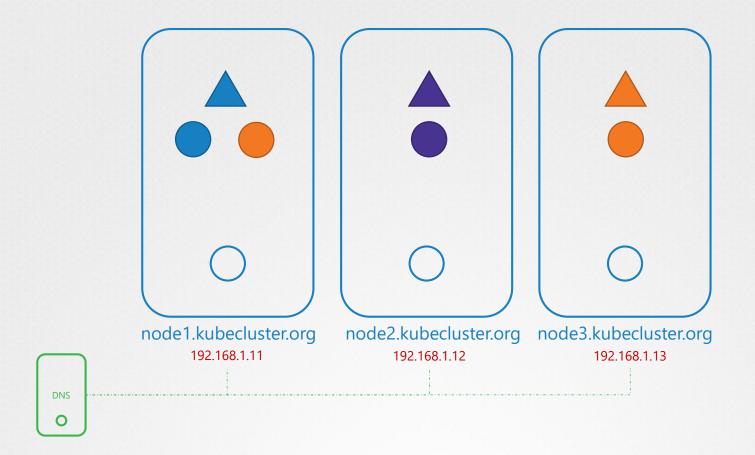
Cluster DNS

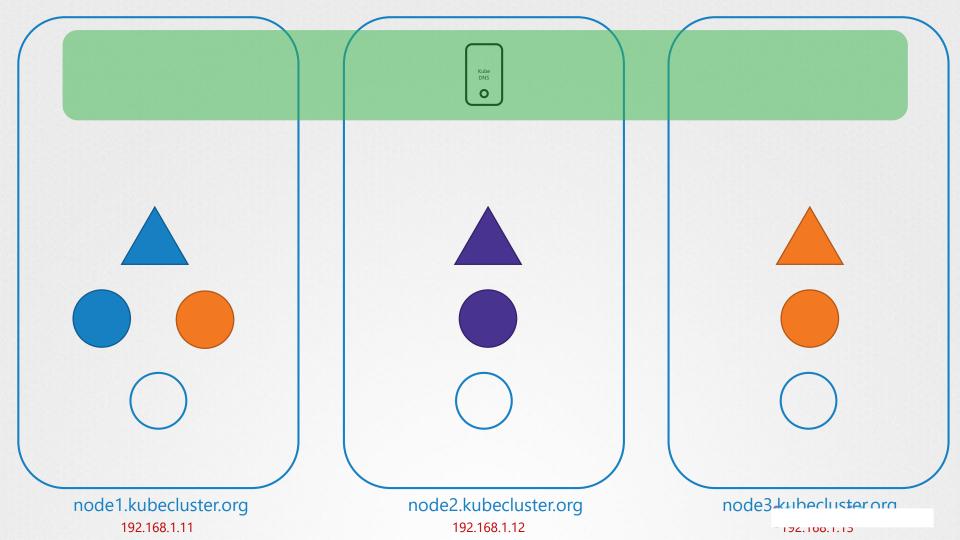
IPre-Requisite

- ✓ What is DNS?
- ✓ Host/NS Lookup, Dig utility
- ✓ Recorded Types A, CNAME
- ✓ Domain Name Hierarchy

IObjectives

- What names are assigned to what objects?
- ☐ Service DNS records
- POD DNS Recrods







| Hostname | IP Address |
|-------------|---------------|
| web-service | 10.107.37.188 |
| | |

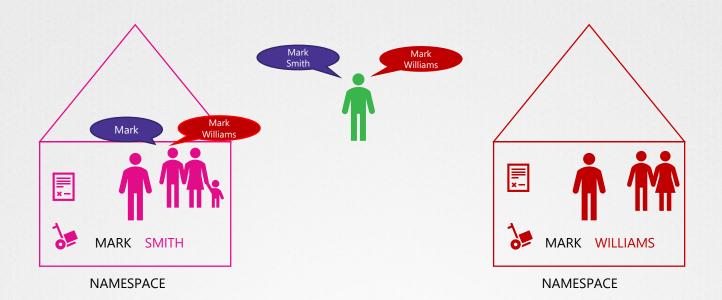


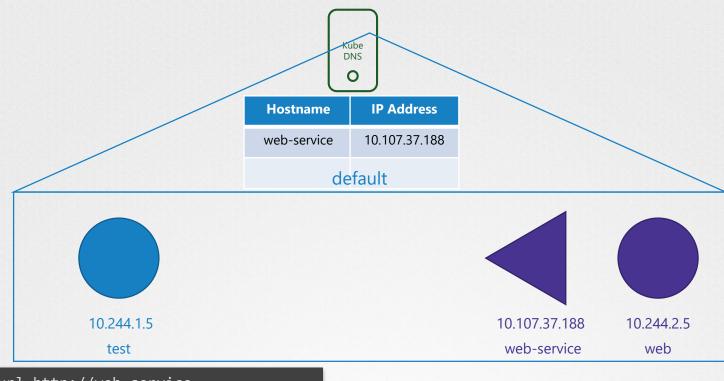
10.244.1.5

test

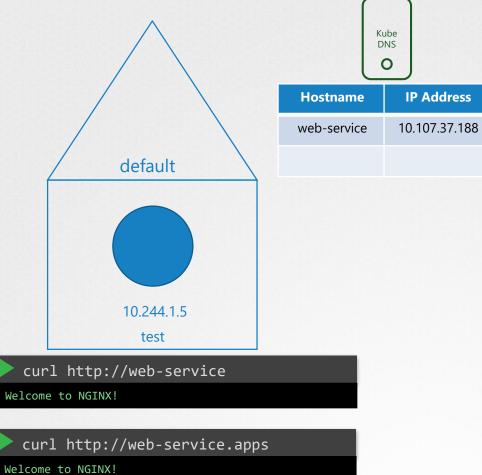


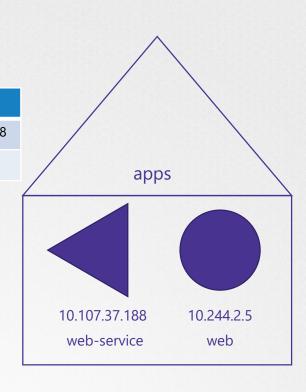
curl http://web-service





curl http://web-service





Kube DNS

| Hostname | Namespace | Туре | Root | IP Address |
|-------------|-----------|------|---------------|---------------|
| web-service | apps | SVC | cluster.local | 10.107.37.188 |

10.107.37.188 10.244.2.5 web-service web

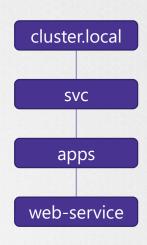
curl http://web-service.apps

Welcome to NGINX!

curl http://web-service.apps.svc

Welcome to NGINX!

curl http://web-service.apps.svc.cluster.local



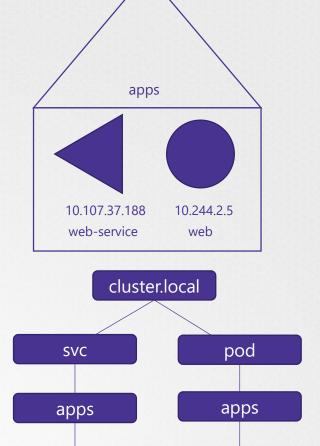


| Hostname | Namespace | Туре | Root | IP Address |
|-------------|-----------|------|---------------|---------------|
| web-service | apps | SVC | cluster.local | 10.107.37.188 |
| 10-244-2-5 | apps | pod | cluster.local | 10.244.2.5 |

10-244-2-5 - 10.244.2.5

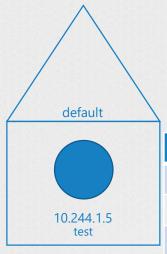
curl http://10-244-2-5.apps.pod.cluster.local

Welcome to NGINX!



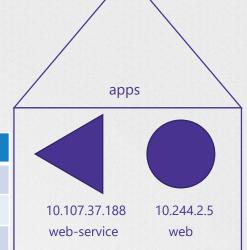
web-service

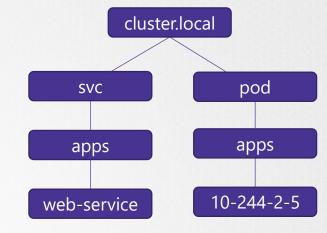
10-244-2-5





| Hostname | Namespace | Туре | Root | IP Address |
|-------------|-----------|------|---------------|---------------|
| web-service | apps | SVC | cluster.local | 10.107.37.188 |
| 10-244-2-5 | apps | pod | cluster.local | 10.244.2.5 |
| 10-244-1-5 | default | pod | cluster.local | 10.244.1.5 |





curl http://10-244-2-5.apps.pod.cluster.local